

modelS_2loss_2optimizers

June 12, 2025

0.1 Importação de bibliotecas

```
[1]: from tensorflow import keras
from keras import layers
from keras.preprocessing import image_dataset_from_directory
import matplotlib.pyplot as plt
from keras.utils import to_categorical
import tensorflow as tf
import numpy as np
from keras.preprocessing import image
from sklearn.metrics import classification_report
import seaborn as sns
import pandas as pd
from sklearn.metrics import confusion_matrix
import os, shutil
```

```
2025-06-12 20:06:21.497301: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1749755181.535213      858 cuda_dnn.cc:8579] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1749755181.543829      858 cuda_blas.cc:1407] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
W0000 00:00:1749755181.575946      858 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
target more than once.
W0000 00:00:1749755181.575973      858 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
target more than once.
W0000 00:00:1749755181.575976      858 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
target more than once.
W0000 00:00:1749755181.575978      858 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
```

target more than once.

2025-06-12 20:06:21.584765: I tensorflow/core/platform/cpu_feature_guard.cc:210]

This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

0.2 Funções

```
[2]: def get_true_pred(model, dataset):  
    y_true = []  
    y_pred = []  
    for images, labels in dataset.unbatch().batch(1):  
        y_true.append(np.argmax(labels.numpy()))  
        pred = model.predict(images, verbose=0)  
        y_pred.append(np.argmax(pred))  
    return np.array(y_true), np.array(y_pred)
```

0.3 Carregamento do dataset

Carrega o dataset distribuido pelos diferentes conjuntos de dados.

```
[3]: train_dir = 'Dataset/archive/seg_train'  
validation_dir = 'Dataset/archive/seg_val'  
test_dir = 'Dataset/archive/seg_test'  
  
train_buildings_dir = 'Dataset/archive/seg_train/buildings/'  
train_forest_dir = 'Dataset/archive/seg_train/forest'  
train_glacier_dir = 'Dataset/archive/seg_train/glacier'  
train_mountain_dir = 'Dataset/archive/seg_train/mountain'  
train_sea_dir = 'Dataset/archive/seg_train/sea'  
train_street_dir = 'Dataset/archive/seg_train/street'  
  
val_buildings_dir = 'Dataset/archive/seg_val/buildings'  
val_forest_dir = 'Dataset/archive/seg_val/forest'  
val_glacier_dir = 'Dataset/archive/seg_val/glacier'  
val_mountain_dir = 'Dataset/archive/seg_val/mountain'  
val_sea_dir = 'Dataset/archive/seg_val/sea'  
val_street_dir = 'Dataset/archive/seg_val/street'  
  
test_buildings_dir = 'Dataset/archive/seg_test/buildings'  
test_forest_dir = 'Dataset/archive/seg_test/forest'  
test_glacier_dir = 'Dataset/archive/seg_test/glacier'  
test_mountain_dir = 'Dataset/archive/seg_test/mountain'  
test_sea_dir = 'Dataset/archive/seg_test/sea'  
test_street_dir = 'Dataset/archive/seg_test/street'  
  
print('total training buildings images:', len(os.listdir(train_buildings_dir)))
```

```

print('total training forest images:', len(os.listdir(train_forest_dir)))
print('total training glacier images:', len(os.listdir(train_glacier_dir)))
print('total training mountain images:', len(os.listdir(train_mountain_dir)))
print('total training sea images:', len(os.listdir(train_sea_dir)))
print('total training street images:', len(os.listdir(train_street_dir)))

print('total validation buildings images:', len(os.listdir(val_buildings_dir)))
print('total validation forest images:', len(os.listdir(val_forest_dir)))
print('total validation glacier images:', len(os.listdir(val_glacier_dir)))
print('total validation mountain images:', len(os.listdir(val_mountain_dir)))
print('total validation sea images:', len(os.listdir(val_sea_dir)))
print('total validation street images:', len(os.listdir(val_street_dir)))

print('total test buildings images:', len(os.listdir(test_buildings_dir)))
print('total test forest images:', len(os.listdir(test_forest_dir)))
print('total test glacier images:', len(os.listdir(test_glacier_dir)))
print('total test mountain images:', len(os.listdir(test_mountain_dir)))
print('total test sea images:', len(os.listdir(test_sea_dir)))
print('total test street images:', len(os.listdir(test_street_dir)))

```

```

total training buildings images: 1691
total training forest images: 1771
total training glacier images: 1904
total training mountain images: 2012
total training sea images: 1774
total training street images: 1882
total validation buildings images: 500
total validation forest images: 500
total validation glacier images: 500
total validation mountain images: 500
total validation sea images: 500
total validation street images: 500
total test buildings images: 437
total test forest images: 474
total test glacier images: 553
total test mountain images: 525
total test sea images: 510
total test street images: 501

```

0.4 Distribuição de imagens por classe e por conjunto de dados

As imagens estão distribuídas por 3 conjuntos de dados: train, validation e test. Cada um desses conjuntos está distribuído por 6 classes: buildings, forest, glacier, mountain, sea e street.

0.4.1 Número total de imagens por classe:

Classe	Treino	Validação	Teste	Total
Buildings	1691	500	437	2628
Forest	1771	500	474	2745
Glacier	1904	500	553	2957
Mountain	2012	500	525	3037
Sea	1774	500	510	2784
Street	1882	500	501	2883
Total	11034	3000	3000	17034

0.4.2 Número total de imagens por conjunto de dados:

Conjunto de dados	Total
Treino	11034
Validação	3000
Teste	3000
Total geral	17034

1 Processamento dos dados

Carrega, redimensiona e organiza imagens em batches com rótulos one-hot, preparando os dados de treino, validação e teste.

```
[4]: IMG_SIZE = 150
    BATCH_SIZE = 32

    # Processing the data
    train_dataset = image_dataset_from_directory(
        train_dir,
        label_mode='categorical',
        image_size=(IMG_SIZE, IMG_SIZE),
        batch_size=BATCH_SIZE)

    validation_dataset = image_dataset_from_directory(
        validation_dir,
        label_mode='categorical',
        image_size=(IMG_SIZE, IMG_SIZE),
        batch_size=BATCH_SIZE)

    test_dataset = image_dataset_from_directory(
        test_dir,
        label_mode='categorical',
        image_size=(IMG_SIZE, IMG_SIZE),
        batch_size=BATCH_SIZE)

    print(test_dataset)
```

```
class_names = train_dataset.class_names
print("Classes:", class_names)
```

Found 11034 files belonging to 6 classes.

```
I0000 00:00:1749755266.582247      858 gpu_device.cc:2019] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 6542 MB memory:  -> device: 0,
name: NVIDIA GeForce GTX 1070, pci bus id: 0000:01:00.0, compute capability: 6.1
```

Found 3000 files belonging to 6 classes.

Found 3000 files belonging to 6 classes.

```
<_PrefetchDataset element_spec=(TensorSpec(shape=(None, 150, 150, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None, 6), dtype=tf.float32,
name=None))>
```

```
Classes: ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

2 Modelo (loss: categorical_crossentropy, optimizer: RMSprop)

2.1 Criação da CNN

Criação da CNN que irá receber imagens de 150x150 píxeis, aplica normalização e passa por quatro camadas convolucionais com max pooling para extrair características, seguidas de uma camada densa com 512 unidades e uma camada de saída softmax para classificação.

```
[18]: inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Flatten()(x)
x = layers.Dense(512, activation="relu")(x)
outputs = layers.Dense(len(class_names), activation="softmax")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

print(model.summary())
```

Model: "functional_2"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 150, 150, 3)	0
rescaling_1 (Rescaling)	(None, 150, 150, 3)	0

conv2d_4 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_5 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_6 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_6 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_7 (Conv2D)	(None, 15, 15, 128)	147,584
max_pooling2d_7 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_2 (Dense)	(None, 512)	3,211,776
dense_3 (Dense)	(None, 6)	3,078

Total params: 3,455,686 (13.18 MB)

Trainable params: 3,455,686 (13.18 MB)

Non-trainable params: 0 (0.00 B)

None

2.2 Compilação da CNN

Compilação da CNN utilizando a loss **categorical_crossentropy** e o optimizer **RMSprop**.

```
[6]: model.compile(
    loss='categorical_crossentropy',
    optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4),
    metrics=['acc'])
```

2.3 Definição do callback

Definição de um callback que guarda automaticamente o modelo com a menor perda (loss) de validação durante o treino.

```
[7]: checkpoint_filepath = 'modelS_CatCross_RMS.keras'
model_checkpoint_callback = keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_filepath,
    monitor='val_loss',
    save_best_only=True)
```

2.4 Treino da CNN

Treino da CNN durante 50 épocas utilizando o dataset de validação e o callback para guardar o melhor modelo.

```
[8]: history_CatCross_RMS = model.fit(
    train_dataset,
    epochs=50,
    validation_data=validation_dataset,
    callbacks=[model_checkpoint_callback])
```

Epoch 1/50

```
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
I0000 00:00:1749755408.539309    1346 service.cc:152] XLA service 0x78e8400049f0
initialized for platform CUDA (this does not guarantee that XLA will be used).
Devices:
I0000 00:00:1749755408.539356    1346 service.cc:160]   StreamExecutor device
(0): NVIDIA GeForce GTX 1070, Compute Capability 6.1
2025-06-12 20:10:08.628419: I
tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR
crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
I0000 00:00:1749755408.886822    1346 cuda_dnn.cc:529] Loaded cuDNN version
90300
2025-06-12 20:10:14.061045: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.12 = (f32[32,32,148,148]{3,2,1,0}, u8[0]{0}) custom-
call(f32[32,3,150,150]{3,2,1,0} %bitcast.2658, f32[32,3,3,3]{3,2,1,0}
%bitcast.2459, f32[32]{0} %bitcast.3030), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":0,"wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}
2025-06-12 20:10:14.120199: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.13 = (f32[32,64,72,72]{3,2,1,0}, u8[0]{0}) custom-
```

```

call(f32[32,32,74,74]{3,2,1,0} %bitcast.3085, f32[64,32,3,3]{3,2,1,0}
%bitcast.2480, f32[64]{0} %bitcast.3125), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1_2/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}
2025-06-12 20:10:14.327237: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.14 = (f32[32,128,34,34]{3,2,1,0}, u8[0]{0}) custom-
call(f32[32,64,36,36]{3,2,1,0} %bitcast.3176, f32[128,64,3,3]{3,2,1,0}
%bitcast.2499, f32[128]{0} %bitcast.3216), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_2_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}
2025-06-12 20:10:14.514173: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.15 = (f32[32,128,15,15]{3,2,1,0}, u8[0]{0}) custom-
call(f32[32,128,17,17]{3,2,1,0} %bitcast.3267, f32[128,128,3,3]{3,2,1,0}
%bitcast.2518, f32[128]{0} %bitcast.3307), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_3_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}

5/345          9s 29ms/step - acc: 0.1741
- loss: 1.7926

I0000 00:00:1749755417.195745    1346 device_compiler.h:188] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.

343/345        0s 30ms/step -
acc: 0.4793 - loss: 1.3048

2025-06-12 20:10:28.002427: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]

```



```

Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.12 = (f32[26,32,148,148]{3,2,1,0}, u8[0]{0}) custom-
call(f32[26,3,150,150]{3,2,1,0} %bitcast.2658, f32[32,3,3,3]{3,2,1,0}
%bitcast.2459, f32[32]{0} %bitcast.3030), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}
2025-06-12 20:10:28.052716: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.13 = (f32[26,64,72,72]{3,2,1,0}, u8[0]{0}) custom-
call(f32[26,32,74,74]{3,2,1,0} %bitcast.3085, f32[64,32,3,3]{3,2,1,0}
%bitcast.2480, f32[64]{0} %bitcast.3125), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1_2/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}
2025-06-12 20:10:28.231026: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.14 = (f32[26,128,34,34]{3,2,1,0}, u8[0]{0}) custom-
call(f32[26,64,36,36]{3,2,1,0} %bitcast.3176, f32[128,64,3,3]{3,2,1,0}
%bitcast.2499, f32[128]{0} %bitcast.3216), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_2_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}
2025-06-12 20:10:28.370095: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.15 = (f32[26,128,15,15]{3,2,1,0}, u8[0]{0}) custom-
call(f32[26,128,17,17]{3,2,1,0} %bitcast.3267, f32[128,128,3,3]{3,2,1,0}
%bitcast.2518, f32[128]{0} %bitcast.3307), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",

```

```
metadata={op_type="Conv2D" op_name="functional_1/conv2d_3_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_conf
ig": {"conv_result_scale": 1, "activation_mode": "kNone", "side_input_scale": 0, "leaky
relu_alpha": 0}, "force_earliest_schedule": false}
```

345/345 0s 39ms/step -
acc: 0.4797 - loss: 1.3040

```
2025-06-12 20:10:31.220247: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.12 = (f32[32,32,148,148]{3,2,1,0}, u8[0]{0}) custom-
call(f32[32,3,150,150]{3,2,1,0} %bitcast.569, f32[32,3,3,3]{3,2,1,0}
%bitcast.576, f32[32]{0} %bitcast.578), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_conf
ig": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leaky
relu_alpha": 0}, "force_earliest_schedule": false}
```

```
2025-06-12 20:10:31.301916: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.13 = (f32[32,64,72,72]{3,2,1,0}, u8[0]{0}) custom-
call(f32[32,32,74,74]{3,2,1,0} %bitcast.586, f32[64,32,3,3]{3,2,1,0}
%bitcast.593, f32[64]{0} %bitcast.595), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1_2/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_conf
ig": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leaky
relu_alpha": 0}, "force_earliest_schedule": false}
```

```
2025-06-12 20:10:31.511840: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.14 = (f32[32,128,34,34]{3,2,1,0}, u8[0]{0}) custom-
call(f32[32,64,36,36]{3,2,1,0} %bitcast.601, f32[128,64,3,3]{3,2,1,0}
%bitcast.608, f32[128]{0} %bitcast.610), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_2_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
```

```

"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 20:10:31.693318: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-activation.15 = (f32[32,128,15,15]{3,2,1,0}, u8[0]{0}) custom-call(f32[32,128,17,17]{3,2,1,0} %bitcast.616, f32[128,128,3,3]{3,2,1,0} %bitcast.623, f32[128]{0} %bitcast.625), window={size=3x3}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn$convBiasActivationForward", metadata={op_type="Conv2D" op_name="functional_1/conv2d_3_1/convolution" source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 20:10:36.686971: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-activation.12 = (f32[24,32,148,148]{3,2,1,0}, u8[0]{0}) custom-call(f32[24,3,150,150]{3,2,1,0} %bitcast.569, f32[32,3,3,3]{3,2,1,0} %bitcast.576, f32[32]{0} %bitcast.578), window={size=3x3}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn$convBiasActivationForward", metadata={op_type="Conv2D" op_name="functional_1/conv2d_1/convolution" source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 20:10:36.732031: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-activation.13 = (f32[24,64,72,72]{3,2,1,0}, u8[0]{0}) custom-call(f32[24,32,74,74]{3,2,1,0} %bitcast.586, f32[64,32,3,3]{3,2,1,0} %bitcast.593, f32[64]{0} %bitcast.595), window={size=3x3}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn$convBiasActivationForward", metadata={op_type="Conv2D" op_name="functional_1/conv2d_1_2/convolution" source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 20:10:36.885420: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-

```

```

activation.14 = (f32[24,128,34,34]{3,2,1,0}, u8[0]{0}) custom-
call(f32[24,64,36,36]{3,2,1,0} %bitcast.601, f32[128,64,3,3]{3,2,1,0}
%bitcast.608, f32[128]{0} %bitcast.610), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_2_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kRelu","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}
2025-06-12 20:10:37.034066: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.15 = (f32[24,128,15,15]{3,2,1,0}, u8[0]{0}) custom-
call(f32[24,128,17,17]{3,2,1,0} %bitcast.616, f32[128,128,3,3]{3,2,1,0}
%bitcast.623, f32[128]{0} %bitcast.625), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_3_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_conf
ig":{"conv_result_scale":1,"activation_mode":"kRelu","side_input_scale":0,"leaky
relu_alpha":0},"force_earliest_schedule":false}

345/345          32s 60ms/step -
acc: 0.4799 - loss: 1.3035 - val_acc: 0.6147 - val_loss: 1.0001
Epoch 2/50
345/345          11s 33ms/step -
acc: 0.6378 - loss: 0.9455 - val_acc: 0.6273 - val_loss: 0.9433
Epoch 3/50
345/345          11s 32ms/step -
acc: 0.6909 - loss: 0.8272 - val_acc: 0.7133 - val_loss: 0.7712
Epoch 4/50
345/345          14s 40ms/step -
acc: 0.7273 - loss: 0.7409 - val_acc: 0.7127 - val_loss: 0.7510
Epoch 5/50
345/345          11s 31ms/step -
acc: 0.7502 - loss: 0.6786 - val_acc: 0.5873 - val_loss: 1.0899
Epoch 6/50
345/345          11s 32ms/step -
acc: 0.7675 - loss: 0.6406 - val_acc: 0.7593 - val_loss: 0.6479
Epoch 7/50
345/345          14s 41ms/step -
acc: 0.7861 - loss: 0.5857 - val_acc: 0.7723 - val_loss: 0.6296
Epoch 8/50
345/345          11s 32ms/step -

```

```

acc: 0.8000 - loss: 0.5504 - val_acc: 0.7933 - val_loss: 0.5771
Epoch 9/50
345/345          11s 32ms/step -
acc: 0.8089 - loss: 0.5090 - val_acc: 0.7880 - val_loss: 0.6025
Epoch 10/50
345/345          14s 41ms/step -
acc: 0.8282 - loss: 0.4808 - val_acc: 0.8070 - val_loss: 0.5311
Epoch 11/50
345/345          11s 31ms/step -
acc: 0.8384 - loss: 0.4425 - val_acc: 0.8063 - val_loss: 0.5388
Epoch 12/50
345/345          11s 32ms/step -
acc: 0.8510 - loss: 0.4095 - val_acc: 0.8173 - val_loss: 0.5204
Epoch 13/50
345/345          14s 41ms/step -
acc: 0.8649 - loss: 0.3892 - val_acc: 0.8210 - val_loss: 0.5094
Epoch 14/50
345/345          11s 33ms/step -
acc: 0.8702 - loss: 0.3635 - val_acc: 0.8250 - val_loss: 0.5011
Epoch 15/50
345/345          11s 31ms/step -
acc: 0.8802 - loss: 0.3297 - val_acc: 0.8277 - val_loss: 0.5056
Epoch 16/50
345/345          14s 40ms/step -
acc: 0.8881 - loss: 0.3096 - val_acc: 0.8220 - val_loss: 0.5235
Epoch 17/50
345/345          11s 31ms/step -
acc: 0.8963 - loss: 0.2870 - val_acc: 0.8223 - val_loss: 0.5169
Epoch 18/50
345/345          11s 31ms/step -
acc: 0.9104 - loss: 0.2554 - val_acc: 0.8317 - val_loss: 0.5126
Epoch 19/50
345/345          14s 40ms/step -
acc: 0.9138 - loss: 0.2410 - val_acc: 0.8317 - val_loss: 0.5212
Epoch 20/50
345/345          11s 32ms/step -
acc: 0.9273 - loss: 0.2118 - val_acc: 0.8260 - val_loss: 0.5310
Epoch 21/50
345/345          11s 32ms/step -
acc: 0.9305 - loss: 0.1972 - val_acc: 0.8223 - val_loss: 0.5691
Epoch 22/50
345/345          14s 40ms/step -
acc: 0.9422 - loss: 0.1724 - val_acc: 0.7290 - val_loss: 1.1331
Epoch 23/50
345/345          11s 31ms/step -
acc: 0.9468 - loss: 0.1567 - val_acc: 0.7937 - val_loss: 0.8158
Epoch 24/50
345/345          11s 31ms/step -

```

acc: 0.9520 - loss: 0.1341 - val_acc: 0.8207 - val_loss: 0.6506
 Epoch 25/50
 345/345 14s 39ms/step -
 acc: 0.9644 - loss: 0.1092 - val_acc: 0.8067 - val_loss: 0.7928
 Epoch 26/50
 345/345 11s 31ms/step -
 acc: 0.9647 - loss: 0.1049 - val_acc: 0.8173 - val_loss: 0.7138
 Epoch 27/50
 345/345 11s 31ms/step -
 acc: 0.9708 - loss: 0.0910 - val_acc: 0.8153 - val_loss: 0.7679
 Epoch 28/50
 345/345 14s 39ms/step -
 acc: 0.9751 - loss: 0.0828 - val_acc: 0.8250 - val_loss: 0.7370
 Epoch 29/50
 345/345 11s 31ms/step -
 acc: 0.9806 - loss: 0.0648 - val_acc: 0.7557 - val_loss: 1.2384
 Epoch 30/50
 345/345 11s 31ms/step -
 acc: 0.9834 - loss: 0.0623 - val_acc: 0.8187 - val_loss: 0.8667
 Epoch 31/50
 345/345 14s 39ms/step -
 acc: 0.9853 - loss: 0.0517 - val_acc: 0.8230 - val_loss: 0.8482
 Epoch 32/50
 345/345 12s 33ms/step -
 acc: 0.9859 - loss: 0.0532 - val_acc: 0.8217 - val_loss: 0.8822
 Epoch 33/50
 345/345 11s 33ms/step -
 acc: 0.9900 - loss: 0.0361 - val_acc: 0.8210 - val_loss: 0.9564
 Epoch 34/50
 345/345 14s 40ms/step -
 acc: 0.9877 - loss: 0.0398 - val_acc: 0.8127 - val_loss: 0.9933
 Epoch 35/50
 345/345 11s 31ms/step -
 acc: 0.9909 - loss: 0.0311 - val_acc: 0.8177 - val_loss: 0.9931
 Epoch 36/50
 345/345 11s 31ms/step -
 acc: 0.9918 - loss: 0.0280 - val_acc: 0.8157 - val_loss: 1.0464
 Epoch 37/50
 345/345 14s 39ms/step -
 acc: 0.9910 - loss: 0.0309 - val_acc: 0.8250 - val_loss: 0.9815
 Epoch 38/50
 345/345 11s 31ms/step -
 acc: 0.9923 - loss: 0.0244 - val_acc: 0.8230 - val_loss: 1.0329
 Epoch 39/50
 345/345 11s 31ms/step -
 acc: 0.9898 - loss: 0.0302 - val_acc: 0.8093 - val_loss: 1.1644
 Epoch 40/50
 345/345 14s 39ms/step -

```

acc: 0.9935 - loss: 0.0317 - val_acc: 0.8297 - val_loss: 1.0730
Epoch 41/50
345/345          11s 31ms/step -
acc: 0.9935 - loss: 0.0218 - val_acc: 0.8017 - val_loss: 1.3321
Epoch 42/50
345/345          11s 31ms/step -
acc: 0.9927 - loss: 0.0280 - val_acc: 0.8277 - val_loss: 1.1037
Epoch 43/50
345/345          14s 39ms/step -
acc: 0.9944 - loss: 0.0182 - val_acc: 0.8173 - val_loss: 1.1486
Epoch 44/50
345/345          11s 31ms/step -
acc: 0.9948 - loss: 0.0158 - val_acc: 0.8313 - val_loss: 1.1271
Epoch 45/50
345/345          11s 31ms/step -
acc: 0.9966 - loss: 0.0117 - val_acc: 0.8237 - val_loss: 1.2193
Epoch 46/50
345/345          14s 39ms/step -
acc: 0.9939 - loss: 0.0221 - val_acc: 0.8010 - val_loss: 1.3957
Epoch 47/50
345/345          11s 31ms/step -
acc: 0.9960 - loss: 0.0113 - val_acc: 0.8170 - val_loss: 1.2861
Epoch 48/50
345/345          11s 31ms/step -
acc: 0.9962 - loss: 0.0149 - val_acc: 0.8213 - val_loss: 1.2919
Epoch 49/50
345/345          14s 39ms/step -
acc: 0.9959 - loss: 0.0138 - val_acc: 0.8187 - val_loss: 1.3313
Epoch 50/50
345/345          11s 31ms/step -
acc: 0.9960 - loss: 0.0111 - val_acc: 0.8187 - val_loss: 1.3664

```

```

[9]: best_epoch = np.argmax(history_CatCross_RMS.history['val_loss']) + 1
      print(f"Melhor época (menor val_loss): {best_epoch}")

```

Melhor época (menor val_loss): 14

2.5 Carregamento do modelo e validação

Carregamento e avaliação do modelo através do valor da accuracy.

```

[10]: modelS_CatCross_RMS = keras.models.load_model('modelS_CatCross_RMS.keras')
      val_loss, val_acc = modelS_CatCross_RMS.evaluate(validation_dataset)
      print('val_acc:', val_acc)

```


```


94/94          3s 19ms/step - acc:
0.8345 - loss: 0.4762
val_acc: 0.824999988079071

```

Representação gráfica dos valores da accuracy e da loss ao longo das épocas.

```
[11]: acc = history_CatCross_RMS.history['acc']
val_acc = history_CatCross_RMS.history['val_acc']
loss = history_CatCross_RMS.history['loss']
val_loss = history_CatCross_RMS.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```





Avaliação da performance do modelo no conjunto de teste, utilizando o relatório de classificação. O relatório apresenta, para cada classe, as métricas precision, recall e F1-score, permitindo analisar detalhadamente os acertos e erros por classe.

```
[12]: y_true, y_pred = get_true_pred(modelS_CatCross_RMS, test_dataset)
report = classification_report(y_true, y_pred, target_names=class_names,
                                output_dict=True)
class_only_report = {k: v for k, v in report.items() if k in class_names}
df = pd.DataFrame(class_only_report).T
print(df[['precision', 'recall', 'f1-score']].round(3))
```

```
2025-06-12 20:22:12.173160: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.12 = (f32[1,32,148,148]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,3,150,150]{3,2,1,0} %bitcast.262, f32[32,3,3,3]{3,2,1,0}
%bitcast.269, f32[32]{0} %bitcast.271), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
```

```

"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 20:22:12.250748: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.13 = (f32[1,64,72,72]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,32,74,74]{3,2,1,0} %bitcast.278, f32[64,32,3,3]{3,2,1,0}
%bitcast.285, f32[64]{0} %bitcast.287), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_1_2/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 20:22:12.354361: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.14 = (f32[1,128,34,34]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,64,36,36]{3,2,1,0} %bitcast.293, f32[128,64,3,3]{3,2,1,0}
%bitcast.300, f32[128]{0} %bitcast.302), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_2_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 20:22:12.451257: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.15 = (f32[1,128,15,15]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,128,17,17]{3,2,1,0} %bitcast.308, f32[128,128,3,3]{3,2,1,0}
%bitcast.315, f32[128]{0} %bitcast.317), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1/conv2d_3_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}

```

	precision	recall	f1-score
buildings	0.778	0.828	0.803

forest	0.959	0.930	0.944
glacier	0.846	0.714	0.775
mountain	0.798	0.792	0.795
sea	0.809	0.849	0.829
street	0.803	0.884	0.841

2025-06-12 20:25:36.900883: I tensorflow/core/framework/local_rendezvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

3 Modelo (loss: KLDivergence, optimizer: SGD)

3.1 Criação da CNN

Criação da CNN que irá receber imagens de 150x150 píxeis, aplica normalização e passa por quatro camadas convolucionais com max pooling para extrair características, seguidas de uma camada densa com 512 unidades e uma camada de saída softmax para classificação.

```
[19]: inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Flatten()(x)
x = layers.Dense(512, activation="relu")(x)
outputs = layers.Dense(len(class_names), activation="softmax")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

print(model.summary())
```

Model: "functional_3"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 150, 150, 3)	0
rescaling_1 (Rescaling)	(None, 150, 150, 3)	0
conv2d_4 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_5 (Conv2D)	(None, 72, 72, 64)	18,496

max_pooling2d_5 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_6 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_6 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_7 (Conv2D)	(None, 15, 15, 128)	147,584
max_pooling2d_7 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_2 (Dense)	(None, 512)	3,211,776
dense_3 (Dense)	(None, 6)	3,078

Total params: 3,455,686 (13.18 MB)

Trainable params: 3,455,686 (13.18 MB)

Non-trainable params: 0 (0.00 B)

None

3.2 Compilação da CNN

Compilação da CNN utilizando a loss **KLDivergence** e o optimizer **SGD**.

```
[22]: model.compile(
      loss=tf.keras.losses.KLDivergence(),
      optimizer=tf.keras.optimizers.SGD(learning_rate=0.01),
      metrics=['acc'])
```

3.3 Definição do callback

Definição de um callback que guarda automaticamente o modelo com a menor perda (loss) de validação durante o treino.

```
[23]: checkpoint_filepath = 'modelS_KLD_SGD.keras'
      model_checkpoint_callback = keras.callbacks.ModelCheckpoint(
          filepath=checkpoint_filepath,
          monitor='val_loss',
          save_best_only=True)
```

3.4 Treino da CNN

Treino da CNN durante 50 épocas utilizando o dataset de validação e o callback para guardar o melhor modelo.

```
[24]: history_KLD_SGD = model.fit(  
    train_dataset,  
    epochs=50,  
    validation_data=validation_dataset,  
    callbacks=[model_checkpoint_callback])
```

```
Epoch 1/50  
345/345          23s 58ms/step -  
acc: 0.2734 - loss: 1.6998 - val_acc: 0.5510 - val_loss: 1.1666  
Epoch 2/50  
345/345          13s 36ms/step -  
acc: 0.5540 - loss: 1.1383 - val_acc: 0.5503 - val_loss: 1.1265  
Epoch 3/50  
345/345          18s 51ms/step -  
acc: 0.5951 - loss: 1.0223 - val_acc: 0.5517 - val_loss: 1.1783  
Epoch 4/50  
345/345          13s 36ms/step -  
acc: 0.6254 - loss: 0.9684 - val_acc: 0.6033 - val_loss: 0.9922  
Epoch 5/50  
345/345          18s 51ms/step -  
acc: 0.6522 - loss: 0.9054 - val_acc: 0.6173 - val_loss: 0.9913  
Epoch 6/50  
345/345          13s 36ms/step -  
acc: 0.6746 - loss: 0.8651 - val_acc: 0.6667 - val_loss: 0.8426  
Epoch 7/50  
345/345          15s 43ms/step -  
acc: 0.6883 - loss: 0.8153 - val_acc: 0.6953 - val_loss: 0.7875  
Epoch 8/50  
345/345          15s 45ms/step -  
acc: 0.7067 - loss: 0.7811 - val_acc: 0.7007 - val_loss: 0.7920  
Epoch 9/50  
345/345          15s 42ms/step -  
acc: 0.7281 - loss: 0.7337 - val_acc: 0.7297 - val_loss: 0.7190  
Epoch 10/50  
345/345          16s 45ms/step -  
acc: 0.7335 - loss: 0.7063 - val_acc: 0.6907 - val_loss: 0.7914  
Epoch 11/50  
345/345          12s 36ms/step -  
acc: 0.7546 - loss: 0.6710 - val_acc: 0.7057 - val_loss: 0.7903  
Epoch 12/50  
345/345          15s 42ms/step -  
acc: 0.7693 - loss: 0.6299 - val_acc: 0.7143 - val_loss: 0.7497  
Epoch 13/50  
345/345          15s 45ms/step -
```

acc: 0.7765 - loss: 0.6087 - val_acc: 0.7107 - val_loss: 0.8038
 Epoch 14/50
 345/345 15s 43ms/step -
 acc: 0.7878 - loss: 0.5744 - val_acc: 0.7243 - val_loss: 0.7106
 Epoch 15/50
 345/345 15s 45ms/step -
 acc: 0.8027 - loss: 0.5369 - val_acc: 0.7737 - val_loss: 0.6126
 Epoch 16/50
 345/345 15s 43ms/step -
 acc: 0.8192 - loss: 0.5055 - val_acc: 0.7527 - val_loss: 0.6946
 Epoch 17/50
 345/345 15s 44ms/step -
 acc: 0.8295 - loss: 0.4680 - val_acc: 0.7340 - val_loss: 0.7086
 Epoch 18/50
 345/345 15s 43ms/step -
 acc: 0.8392 - loss: 0.4406 - val_acc: 0.7810 - val_loss: 0.5993
 Epoch 19/50
 345/345 12s 36ms/step -
 acc: 0.8525 - loss: 0.4081 - val_acc: 0.7920 - val_loss: 0.6009
 Epoch 20/50
 345/345 18s 51ms/step -
 acc: 0.8643 - loss: 0.3793 - val_acc: 0.7820 - val_loss: 0.6240
 Epoch 21/50
 345/345 13s 36ms/step -
 acc: 0.8805 - loss: 0.3434 - val_acc: 0.7807 - val_loss: 0.6528
 Epoch 22/50
 345/345 18s 52ms/step -
 acc: 0.8928 - loss: 0.3116 - val_acc: 0.7893 - val_loss: 0.6325
 Epoch 23/50
 345/345 13s 36ms/step -
 acc: 0.9108 - loss: 0.2693 - val_acc: 0.7313 - val_loss: 0.8749
 Epoch 24/50
 345/345 16s 46ms/step -
 acc: 0.9153 - loss: 0.2499 - val_acc: 0.7850 - val_loss: 0.6652
 Epoch 25/50
 345/345 14s 41ms/step -
 acc: 0.9293 - loss: 0.2226 - val_acc: 0.7927 - val_loss: 0.6933
 Epoch 26/50
 345/345 12s 36ms/step -
 acc: 0.9388 - loss: 0.1840 - val_acc: 0.7593 - val_loss: 0.8301
 Epoch 27/50
 345/345 18s 52ms/step -
 acc: 0.9419 - loss: 0.1657 - val_acc: 0.7917 - val_loss: 0.7645
 Epoch 28/50
 345/345 12s 36ms/step -
 acc: 0.9375 - loss: 0.1975 - val_acc: 0.7853 - val_loss: 0.8055
 Epoch 29/50
 345/345 18s 51ms/step -

acc: 0.9615 - loss: 0.1263 - val_acc: 0.7900 - val_loss: 0.7972
 Epoch 30/50
 345/345 12s 36ms/step -
 acc: 0.9710 - loss: 0.0907 - val_acc: 0.7877 - val_loss: 0.8953
 Epoch 31/50
 345/345 18s 52ms/step -
 acc: 0.9648 - loss: 0.1182 - val_acc: 0.7757 - val_loss: 0.9071
 Epoch 32/50
 345/345 12s 36ms/step -
 acc: 0.9802 - loss: 0.0796 - val_acc: 0.7840 - val_loss: 0.9276
 Epoch 33/50
 345/345 14s 41ms/step -
 acc: 0.9790 - loss: 0.0851 - val_acc: 0.7763 - val_loss: 0.9262
 Epoch 34/50
 345/345 16s 46ms/step -
 acc: 0.9784 - loss: 0.0739 - val_acc: 0.7877 - val_loss: 0.9464
 Epoch 35/50
 345/345 12s 36ms/step -
 acc: 0.9870 - loss: 0.0508 - val_acc: 0.7957 - val_loss: 0.9720
 Epoch 36/50
 345/345 18s 53ms/step -
 acc: 0.9922 - loss: 0.0327 - val_acc: 0.7940 - val_loss: 1.0449
 Epoch 37/50
 345/345 13s 37ms/step -
 acc: 0.9924 - loss: 0.0298 - val_acc: 0.7920 - val_loss: 1.0758
 Epoch 38/50
 345/345 15s 42ms/step -
 acc: 0.9943 - loss: 0.0253 - val_acc: 0.7917 - val_loss: 1.0968
 Epoch 39/50
 345/345 15s 45ms/step -
 acc: 0.9604 - loss: 0.2109 - val_acc: 0.7860 - val_loss: 1.0312
 Epoch 40/50
 345/345 14s 42ms/step -
 acc: 0.9909 - loss: 0.0437 - val_acc: 0.7893 - val_loss: 1.0458
 Epoch 41/50
 345/345 15s 44ms/step -
 acc: 0.9971 - loss: 0.0200 - val_acc: 0.7873 - val_loss: 1.1056
 Epoch 42/50
 345/345 15s 43ms/step -
 acc: 0.9943 - loss: 0.0246 - val_acc: 0.7837 - val_loss: 1.1357
 Epoch 43/50
 345/345 15s 44ms/step -
 acc: 0.9780 - loss: 0.1248 - val_acc: 0.7567 - val_loss: 1.2022
 Epoch 44/50
 345/345 15s 43ms/step -
 acc: 0.9897 - loss: 0.0439 - val_acc: 0.7820 - val_loss: 1.0372
 Epoch 45/50
 345/345 13s 39ms/step -

```

acc: 0.9882 - loss: 0.0417 - val_acc: 0.7960 - val_loss: 1.0720
Epoch 46/50
345/345          18s 51ms/step -
acc: 0.9976 - loss: 0.0154 - val_acc: 0.7963 - val_loss: 1.1339
Epoch 47/50
345/345          15s 43ms/step -
acc: 0.9970 - loss: 0.0145 - val_acc: 0.7950 - val_loss: 1.1477
Epoch 48/50
345/345          15s 45ms/step -
acc: 0.9983 - loss: 0.0121 - val_acc: 0.7913 - val_loss: 1.1862
Epoch 49/50
345/345          15s 44ms/step -
acc: 0.9988 - loss: 0.0109 - val_acc: 0.7980 - val_loss: 1.1894
Epoch 50/50
345/345          15s 45ms/step -
acc: 0.9981 - loss: 0.0100 - val_acc: 0.7970 - val_loss: 1.2049

```

```

[25]: best_epoch = np.argmin(history_KLD_SGD.history['val_loss']) + 1
      print(f"Melhor época (menor val_loss): {best_epoch}")

```

Melhor época (menor val_loss): 18

3.5 Carregamento do modelo e validação

Carregamento e avaliação do modelo através do valor da accuracy.

```

[26]: modelS_KLD_SGD = keras.models.load_model('modelS_KLD_SGD.keras')
      val_loss, val_acc = modelS_KLD_SGD.evaluate(validation_dataset)
      print('val_acc:', val_acc)

```

```

94/94          4s 28ms/step - acc:
0.7871 - loss: 0.5720
val_acc: 0.781000018119812

```

Representação gráfica dos valores da accuracy e da loss ao longo das épocas.


```


[27]: acc = history_KLD_SGD.history['acc']
      val_acc = history_KLD_SGD.history['val_acc']
      loss = history_KLD_SGD.history['loss']
      val_loss = history_KLD_SGD.history['val_loss']
      epochs = range(1, len(acc) + 1)
      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.legend()
      plt.figure()
      plt.plot(epochs, loss, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and validation loss')

```



```
plt.legend()  
plt.show()
```





Avaliação da performance do modelo no conjunto de teste, utilizando o relatório de classificação. O relatório apresenta, para cada classe, as métricas precision, recall e F1-score, permitindo analisar detalhadamente os acertos e erros por classe.

```

[28]: y_true, y_pred = get_true_pred(modelS_KLD_SGD, test_dataset)
      report = classification_report(y_true, y_pred, target_names=class_names,
      ↪output_dict=True)
      class_only_report = {k: v for k, v in report.items() if k in class_names}
      df = pd.DataFrame(class_only_report).T
      print(df[['precision', 'recall', 'f1-score']].round(3))
  
```

	precision	recall	f1-score
buildings	0.732	0.769	0.750
forest	0.933	0.876	0.903
glacier	0.689	0.807	0.743
mountain	0.774	0.737	0.755
sea	0.771	0.747	0.759
street	0.846	0.768	0.805

```

2025-06-12 20:52:07.769680: I tensorflow/core/framework/local_rendevvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
  
```

4 Model (loss: categorical_crossentropy, optimizer: SGD)

4.1 Criação da CNN

Criação da CNN que irá receber imagens de 150x150 píxeis, aplica normalização e passa por quatro camadas convolucionais com max pooling para extrair características, seguidas de uma camada densa com 512 unidades e uma camada de saída softmax para classificação.

```
[33]: inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Flatten()(x)
x = layers.Dense(512, activation="relu")(x)
outputs = layers.Dense(len(class_names), activation="softmax")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

print(model.summary())
```

Model: "functional_5"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 150, 150, 3)	0
rescaling_2 (Rescaling)	(None, 150, 150, 3)	0
conv2d_8 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_8 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_9 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_9 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_10 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_10 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_11 (Conv2D)	(None, 15, 15, 128)	147,584

max_pooling2d_11 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_2 (Flatten)	(None, 6272)	0
dense_4 (Dense)	(None, 512)	3,211,776
dense_5 (Dense)	(None, 6)	3,078

Total params: 3,455,686 (13.18 MB)

Trainable params: 3,455,686 (13.18 MB)

Non-trainable params: 0 (0.00 B)

None

4.2 Compilação da CNN

Compilação da CNN utilizando a loss **categorical_crossentropy** e o optimizer **SGD**.

```
[34]: model.compile(
    loss='categorical_crossentropy',
    optimizer=tf.keras.optimizers.SGD(learning_rate=0.01),
    metrics=['acc'])
```

4.3 Definição do callback

Definição de um callback que guarda automaticamente o modelo com a menor perda (loss) de validação durante o treino.

```
[35]: checkpoint_filepath = 'modelS_CatCross_SGD.keras'
model_checkpoint_callback = keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_filepath,
    monitor='val_loss',
    save_best_only=True)
```

4.4 Treino da CNN

Treino da CNN durante 50 épocas utilizando o dataset de validação e o callback para guardar o melhor modelo.

```
[36]: history_CatCross_SGD = model.fit(
    train_dataset,
    epochs=50,
    validation_data=validation_dataset,
    callbacks=[model_checkpoint_callback])
```

Epoch 1/50
 345/345 17s 42ms/step -
 acc: 0.3117 - loss: 1.6749 - val_acc: 0.5253 - val_loss: 1.2196

Epoch 2/50
 345/345 16s 47ms/step -
 acc: 0.5506 - loss: 1.1587 - val_acc: 0.5393 - val_loss: 1.1288

Epoch 3/50
 345/345 12s 35ms/step -
 acc: 0.6010 - loss: 1.0195 - val_acc: 0.6223 - val_loss: 0.9644

Epoch 4/50
 345/345 16s 47ms/step -
 acc: 0.6281 - loss: 0.9593 - val_acc: 0.6510 - val_loss: 0.9041

Epoch 5/50
 345/345 14s 40ms/step -
 acc: 0.6553 - loss: 0.8868 - val_acc: 0.6363 - val_loss: 0.9463

Epoch 6/50
 345/345 12s 36ms/step -
 acc: 0.6709 - loss: 0.8480 - val_acc: 0.7040 - val_loss: 0.7910

Epoch 7/50
 345/345 15s 44ms/step -
 acc: 0.7075 - loss: 0.7865 - val_acc: 0.7247 - val_loss: 0.7414

Epoch 8/50
 345/345 12s 34ms/step -
 acc: 0.7215 - loss: 0.7476 - val_acc: 0.7040 - val_loss: 0.7952

Epoch 9/50
 345/345 15s 42ms/step -
 acc: 0.7356 - loss: 0.7172 - val_acc: 0.6243 - val_loss: 1.0020

Epoch 10/50
 345/345 12s 35ms/step -
 acc: 0.7454 - loss: 0.6819 - val_acc: 0.7433 - val_loss: 0.6923

Epoch 11/50
 345/345 13s 38ms/step -
 acc: 0.7622 - loss: 0.6444 - val_acc: 0.6940 - val_loss: 0.8103

Epoch 12/50
 345/345 17s 49ms/step -
 acc: 0.7810 - loss: 0.6091 - val_acc: 0.5783 - val_loss: 1.1887

Epoch 13/50
 345/345 13s 36ms/step -
 acc: 0.7852 - loss: 0.5973 - val_acc: 0.7653 - val_loss: 0.6421

Epoch 14/50
 345/345 15s 44ms/step -
 acc: 0.8089 - loss: 0.5397 - val_acc: 0.7817 - val_loss: 0.6116

Epoch 15/50
 345/345 12s 35ms/step -
 acc: 0.8095 - loss: 0.5182 - val_acc: 0.7807 - val_loss: 0.6194

Epoch 16/50
 345/345 12s 34ms/step -
 acc: 0.8243 - loss: 0.4847 - val_acc: 0.7887 - val_loss: 0.6118

Epoch 17/50
 345/345 15s 43ms/step -
 acc: 0.8380 - loss: 0.4542 - val_acc: 0.7667 - val_loss: 0.6719
 Epoch 18/50
 345/345 12s 36ms/step -
 acc: 0.8470 - loss: 0.4336 - val_acc: 0.7920 - val_loss: 0.5944
 Epoch 19/50
 345/345 12s 35ms/step -
 acc: 0.8610 - loss: 0.3895 - val_acc: 0.7623 - val_loss: 0.6811
 Epoch 20/50
 345/345 15s 43ms/step -
 acc: 0.8684 - loss: 0.3711 - val_acc: 0.7850 - val_loss: 0.6264
 Epoch 21/50
 345/345 12s 34ms/step -
 acc: 0.8787 - loss: 0.3387 - val_acc: 0.7473 - val_loss: 0.7768
 Epoch 22/50
 345/345 15s 43ms/step -
 acc: 0.8974 - loss: 0.3051 - val_acc: 0.7980 - val_loss: 0.6321
 Epoch 23/50
 345/345 12s 34ms/step -
 acc: 0.9005 - loss: 0.2830 - val_acc: 0.7863 - val_loss: 0.6889
 Epoch 24/50
 345/345 12s 34ms/step -
 acc: 0.9069 - loss: 0.2665 - val_acc: 0.7953 - val_loss: 0.6613
 Epoch 25/50
 345/345 15s 43ms/step -
 acc: 0.9285 - loss: 0.2112 - val_acc: 0.5633 - val_loss: 1.6123
 Epoch 26/50
 345/345 12s 34ms/step -
 acc: 0.9263 - loss: 0.2168 - val_acc: 0.7997 - val_loss: 0.7278
 Epoch 27/50
 345/345 12s 34ms/step -
 acc: 0.9488 - loss: 0.1642 - val_acc: 0.7840 - val_loss: 0.7833
 Epoch 28/50
 345/345 15s 43ms/step -
 acc: 0.9584 - loss: 0.1344 - val_acc: 0.8060 - val_loss: 0.7441
 Epoch 29/50
 345/345 12s 35ms/step -
 acc: 0.9481 - loss: 0.1797 - val_acc: 0.7923 - val_loss: 0.8098
 Epoch 30/50
 345/345 12s 34ms/step -
 acc: 0.9699 - loss: 0.1015 - val_acc: 0.8057 - val_loss: 0.8770
 Epoch 31/50
 345/345 15s 43ms/step -
 acc: 0.9738 - loss: 0.0851 - val_acc: 0.7990 - val_loss: 0.8412
 Epoch 32/50
 345/345 11s 31ms/step -
 acc: 0.9811 - loss: 0.0646 - val_acc: 0.7650 - val_loss: 1.0980

Epoch 33/50
 345/345 11s 30ms/step -
 acc: 0.9657 - loss: 0.1248 - val_acc: 0.8043 - val_loss: 0.8360
 Epoch 34/50
 345/345 14s 39ms/step -
 acc: 0.9856 - loss: 0.0554 - val_acc: 0.7907 - val_loss: 0.9453
 Epoch 35/50
 345/345 11s 31ms/step -
 acc: 0.9873 - loss: 0.0477 - val_acc: 0.7973 - val_loss: 0.9590
 Epoch 36/50
 345/345 11s 30ms/step -
 acc: 0.9885 - loss: 0.0419 - val_acc: 0.8030 - val_loss: 0.9738
 Epoch 37/50
 345/345 13s 39ms/step -
 acc: 0.9863 - loss: 0.0637 - val_acc: 0.7987 - val_loss: 0.9227
 Epoch 38/50
 345/345 11s 30ms/step -
 acc: 0.9953 - loss: 0.0250 - val_acc: 0.8057 - val_loss: 1.0174
 Epoch 39/50
 345/345 11s 31ms/step -
 acc: 0.9953 - loss: 0.0202 - val_acc: 0.8090 - val_loss: 1.0289
 Epoch 40/50
 345/345 13s 39ms/step -
 acc: 0.9785 - loss: 0.0870 - val_acc: 0.8010 - val_loss: 1.0142
 Epoch 41/50
 345/345 11s 30ms/step -
 acc: 0.9916 - loss: 0.0374 - val_acc: 0.8020 - val_loss: 0.9467
 Epoch 42/50
 345/345 11s 30ms/step -
 acc: 0.9979 - loss: 0.0161 - val_acc: 0.8073 - val_loss: 1.0257
 Epoch 43/50
 345/345 13s 39ms/step -
 acc: 0.9976 - loss: 0.0143 - val_acc: 0.8087 - val_loss: 1.0781
 Epoch 44/50
 345/345 11s 30ms/step -
 acc: 0.9972 - loss: 0.0159 - val_acc: 0.8080 - val_loss: 1.1023
 Epoch 45/50
 345/345 10s 30ms/step -
 acc: 0.9976 - loss: 0.0123 - val_acc: 0.8043 - val_loss: 1.1398
 Epoch 46/50
 345/345 13s 39ms/step -
 acc: 0.9970 - loss: 0.0131 - val_acc: 0.8070 - val_loss: 1.1397
 Epoch 47/50
 345/345 11s 30ms/step -
 acc: 0.9849 - loss: 0.0760 - val_acc: 0.8107 - val_loss: 1.0579
 Epoch 48/50
 345/345 11s 30ms/step -
 acc: 0.9978 - loss: 0.0128 - val_acc: 0.8020 - val_loss: 1.1259

```
Epoch 49/50
345/345          13s 39ms/step -
acc: 0.9922 - loss: 0.0393 - val_acc: 0.8040 - val_loss: 1.1032
Epoch 50/50
345/345          11s 31ms/step -
acc: 0.9950 - loss: 0.0205 - val_acc: 0.8083 - val_loss: 1.1324
```

```
[37]: best_epoch = np.argmin(history_CatCross_SGD.history['val_loss']) + 1
      print(f"Melhor época (menor val_loss): {best_epoch}")
```

Melhor época (menor val_loss): 18

4.5 Carregamento do modelo e validação


Carregamento e avaliação do modelo através do valor da accuracy.


```
[38]: modelS_CatCross_SGD = keras.models.load_model('modelS_CatCross_SGD.keras')
      val_loss, val_acc = modelS_CatCross_SGD.evaluate(validation_dataset)
      print('val_acc:', val_acc)
```

```
94/94          3s 19ms/step - acc:
0.7904 - loss: 0.5882
val_acc: 0.7919999957084656
```

Representação gráfica dos valores da accuracy e da loss ao longo das épocas.

```
[39]: acc = history_CatCross_SGD.history['acc']
      val_acc = history_CatCross_SGD.history['val_acc']
      loss = history_CatCross_SGD.history['loss']
      val_loss = history_CatCross_SGD.history['val_loss']
      epochs = range(1, len(acc) + 1)
      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.legend()
      plt.figure()
      plt.plot(epochs, loss, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and validation loss')
      plt.legend()
      plt.show()
```



Avaliação da performance do modelo no conjunto de teste, utilizando o relatório de classificação. O relatório apresenta, para cada classe, as métricas precision, recall e F1-score, permitindo analisar detalhadamente os acertos e erros por classe.

```
[40]: y_true, y_pred = get_true_pred(modelS_CatCross_SGD, test_dataset)
report = classification_report(y_true, y_pred, target_names=class_names,
                               output_dict=True)
class_only_report = {k: v for k, v in report.items() if k in class_names}
df = pd.DataFrame(class_only_report).T
print(df[['precision', 'recall', 'f1-score']].round(3))
```

	precision	recall	f1-score
buildings	0.773	0.773	0.773
forest	0.883	0.937	0.909
glacier	0.689	0.826	0.752
mountain	0.777	0.758	0.768
sea	0.833	0.731	0.779
street	0.872	0.760	0.812

5 Model (loss: KLDivergence, optimizer: RMSprop)

5.1 Criação da CNN

Criação da CNN que irá receber imagens de 150x150 píxeis, aplica normalização e passa por quatro camadas convolucionais com max pooling para extrair características, seguidas de uma camada densa com 512 unidades e uma camada de saída softmax para classificação.

```
[54]: inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Flatten()(x)
x = layers.Dense(512, activation="relu")(x)
outputs = layers.Dense(len(class_names), activation="softmax")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

print(model.summary())
```

Model: "functional_7"

Layer (type)	Output Shape	Param #
input_layer_4 (InputLayer)	(None, 150, 150, 3)	0
rescaling_4 (Rescaling)	(None, 150, 150, 3)	0
conv2d_16 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_16 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_17 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_17 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_18 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_18 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_19 (Conv2D)	(None, 15, 15, 128)	147,584

max_pooling2d_19 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_4 (Flatten)	(None, 6272)	0
dense_8 (Dense)	(None, 512)	3,211,776
dense_9 (Dense)	(None, 6)	3,078

Total params: 3,455,686 (13.18 MB)

Trainable params: 3,455,686 (13.18 MB)

Non-trainable params: 0 (0.00 B)

None

5.2 Compilação da CNN

Compilação da CNN utilizando a loss **KLDivergence** e o optimizer **RMSprop**.

```
[55]: model.compile(
        loss=tf.keras.losses.KLDivergence(),
        optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4),
        metrics=['acc'])
```

5.3 Definição do callback

Definição de um callback que guarda automaticamente o modelo com a menor perda (loss) de validação durante o treino.

```
[56]: checkpoint_filepath = 'modelS_KLD_RMS.keras'
model_checkpoint_callback = keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_filepath,
    monitor='val_loss',
    save_best_only=True)
```

5.4 Treino da CNN

Treino da CNN durante 50 épocas utilizando o dataset de validação e o callback para guardar o melhor modelo.

```
[57]: history_KLD_RMS = model.fit(
    train_dataset,
    epochs=50,
    validation_data=validation_dataset,
    callbacks=[model_checkpoint_callback])
```

Epoch 1/50
 345/345 20s 48ms/step -
 acc: 0.4663 - loss: 1.3290 - val_acc: 0.5753 - val_loss: 1.0506
 Epoch 2/50
 345/345 12s 34ms/step -
 acc: 0.6347 - loss: 0.9598 - val_acc: 0.6423 - val_loss: 0.9260
 Epoch 3/50
 345/345 12s 34ms/step -
 acc: 0.6871 - loss: 0.8445 - val_acc: 0.7303 - val_loss: 0.7340
 Epoch 4/50
 345/345 15s 42ms/step -
 acc: 0.7180 - loss: 0.7534 - val_acc: 0.6693 - val_loss: 0.9300
 Epoch 5/50
 345/345 12s 35ms/step -
 acc: 0.7488 - loss: 0.6864 - val_acc: 0.7473 - val_loss: 0.6989
 Epoch 6/50
 345/345 12s 35ms/step -
 acc: 0.7696 - loss: 0.6371 - val_acc: 0.7650 - val_loss: 0.6384
 Epoch 7/50
 345/345 15s 43ms/step -
 acc: 0.7939 - loss: 0.5875 - val_acc: 0.7540 - val_loss: 0.6796
 Epoch 8/50
 345/345 12s 34ms/step -
 acc: 0.8059 - loss: 0.5519 - val_acc: 0.7943 - val_loss: 0.5728
 Epoch 9/50
 345/345 12s 34ms/step -
 acc: 0.8168 - loss: 0.5104 - val_acc: 0.7910 - val_loss: 0.5759
 Epoch 10/50
 345/345 15s 43ms/step -
 acc: 0.8259 - loss: 0.4722 - val_acc: 0.8153 - val_loss: 0.5304
 Epoch 11/50
 345/345 12s 34ms/step -
 acc: 0.8423 - loss: 0.4431 - val_acc: 0.8190 - val_loss: 0.5000
 Epoch 12/50
 345/345 15s 42ms/step -
 acc: 0.8492 - loss: 0.4153 - val_acc: 0.7613 - val_loss: 0.6571
 Epoch 13/50
 345/345 12s 34ms/step -
 acc: 0.8557 - loss: 0.3983 - val_acc: 0.8190 - val_loss: 0.4973
 Epoch 14/50
 345/345 12s 34ms/step -
 acc: 0.8693 - loss: 0.3681 - val_acc: 0.8120 - val_loss: 0.5250
 Epoch 15/50
 345/345 15s 42ms/step -
 acc: 0.8793 - loss: 0.3397 - val_acc: 0.8170 - val_loss: 0.5077
 Epoch 16/50
 345/345 12s 34ms/step -
 acc: 0.8847 - loss: 0.3175 - val_acc: 0.7830 - val_loss: 0.6482

Epoch 17/50
 345/345 13s 37ms/step -
 acc: 0.8963 - loss: 0.2932 - val_acc: 0.8113 - val_loss: 0.5510
 Epoch 18/50
 345/345 16s 47ms/step -
 acc: 0.9061 - loss: 0.2649 - val_acc: 0.8047 - val_loss: 0.5980
 Epoch 19/50
 345/345 13s 38ms/step -
 acc: 0.9142 - loss: 0.2475 - val_acc: 0.8200 - val_loss: 0.5531
 Epoch 20/50
 345/345 18s 51ms/step -
 acc: 0.9223 - loss: 0.2218 - val_acc: 0.8143 - val_loss: 0.5870
 Epoch 21/50
 345/345 15s 42ms/step -
 acc: 0.9336 - loss: 0.1979 - val_acc: 0.8220 - val_loss: 0.5715
 Epoch 22/50
 345/345 17s 48ms/step -
 acc: 0.9401 - loss: 0.1773 - val_acc: 0.8280 - val_loss: 0.5440
 Epoch 23/50
 345/345 15s 43ms/step -
 acc: 0.9466 - loss: 0.1509 - val_acc: 0.8167 - val_loss: 0.6117
 Epoch 24/50
 345/345 14s 40ms/step -
 acc: 0.9550 - loss: 0.1352 - val_acc: 0.8240 - val_loss: 0.6028
 Epoch 25/50
 345/345 16s 47ms/step -
 acc: 0.9626 - loss: 0.1194 - val_acc: 0.8230 - val_loss: 0.6811
 Epoch 26/50
 345/345 15s 43ms/step -
 acc: 0.9687 - loss: 0.1016 - val_acc: 0.8290 - val_loss: 0.6299
 Epoch 27/50
 345/345 17s 49ms/step -
 acc: 0.9734 - loss: 0.0893 - val_acc: 0.8087 - val_loss: 0.7277
 Epoch 28/50
 345/345 14s 40ms/step -
 acc: 0.9738 - loss: 0.0777 - val_acc: 0.8217 - val_loss: 0.7308
 Epoch 29/50
 345/345 16s 47ms/step -
 acc: 0.9819 - loss: 0.0615 - val_acc: 0.8273 - val_loss: 0.7720
 Epoch 30/50
 345/345 13s 38ms/step -
 acc: 0.9816 - loss: 0.0612 - val_acc: 0.8140 - val_loss: 0.8927
 Epoch 31/50
 345/345 12s 36ms/step -
 acc: 0.9839 - loss: 0.0554 - val_acc: 0.8133 - val_loss: 0.8512
 Epoch 32/50
 345/345 15s 44ms/step -
 acc: 0.9856 - loss: 0.0479 - val_acc: 0.8093 - val_loss: 0.8768

Epoch 33/50
 345/345 12s 36ms/step -
 acc: 0.9887 - loss: 0.0416 - val_acc: 0.8187 - val_loss: 0.8435
 Epoch 34/50
 345/345 16s 47ms/step -
 acc: 0.9900 - loss: 0.0310 - val_acc: 0.8247 - val_loss: 0.8899
 Epoch 35/50
 345/345 13s 36ms/step -
 acc: 0.9912 - loss: 0.0306 - val_acc: 0.7297 - val_loss: 1.7118
 Epoch 36/50
 345/345 12s 36ms/step -
 acc: 0.9886 - loss: 0.0465 - val_acc: 0.8210 - val_loss: 0.9315
 Epoch 37/50
 345/345 15s 44ms/step -
 acc: 0.9898 - loss: 0.0351 - val_acc: 0.8243 - val_loss: 0.9315
 Epoch 38/50
 345/345 12s 36ms/step -
 acc: 0.9931 - loss: 0.0252 - val_acc: 0.8177 - val_loss: 1.0275
 Epoch 39/50
 345/345 15s 44ms/step -
 acc: 0.9919 - loss: 0.0294 - val_acc: 0.8183 - val_loss: 1.0122
 Epoch 40/50
 345/345 12s 35ms/step -
 acc: 0.9929 - loss: 0.0257 - val_acc: 0.8213 - val_loss: 1.0232
 Epoch 41/50
 345/345 12s 36ms/step -
 acc: 0.9936 - loss: 0.0247 - val_acc: 0.8197 - val_loss: 1.0245
 Epoch 42/50
 345/345 15s 44ms/step -
 acc: 0.9934 - loss: 0.0262 - val_acc: 0.8237 - val_loss: 1.0647
 Epoch 43/50
 345/345 12s 35ms/step -
 acc: 0.9951 - loss: 0.0162 - val_acc: 0.8177 - val_loss: 1.0867
 Epoch 44/50
 345/345 12s 36ms/step -
 acc: 0.9949 - loss: 0.0173 - val_acc: 0.8210 - val_loss: 1.1388
 Epoch 45/50
 345/345 15s 44ms/step -
 acc: 0.9948 - loss: 0.0163 - val_acc: 0.8233 - val_loss: 1.0966
 Epoch 46/50
 345/345 12s 35ms/step -
 acc: 0.9949 - loss: 0.0188 - val_acc: 0.8210 - val_loss: 1.1640
 Epoch 47/50
 345/345 16s 46ms/step -
 acc: 0.9931 - loss: 0.0187 - val_acc: 0.8213 - val_loss: 1.1654
 Epoch 48/50
 345/345 13s 38ms/step -
 acc: 0.9951 - loss: 0.0167 - val_acc: 0.8207 - val_loss: 1.1432

```
Epoch 49/50
345/345          12s 36ms/step -
acc: 0.9957 - loss: 0.0167 - val_acc: 0.8247 - val_loss: 1.1050
Epoch 50/50
345/345          17s 48ms/step -
acc: 0.9972 - loss: 0.0120 - val_acc: 0.8170 - val_loss: 1.1817
```

```
[58]: best_epoch = np.argmin(history_KLD_RMS.history['val_loss']) + 1
      print(f"Melhor época (menor val_loss): {best_epoch}")
```

Melhor época (menor val_loss): 13

5.5 Carregamento do modelo e validação


Carregamento e avaliação do modelo através do valor da accuracy.


```
[59]: modelS_KLD_RMS = keras.models.load_model('modelS_KLD_RMS.keras')
      val_loss, val_acc = modelS_KLD_RMS.evaluate(validation_dataset)
      print('val_acc:', val_acc)
```

```
94/94          3s 18ms/step - acc:
0.8298 - loss: 0.4612
val_acc: 0.8190000057220459
```

Representação gráfica dos valores da accuracy e da loss ao longo das épocas.

```
[47]: acc = history_KLD_RMS.history['acc']
      val_acc = history_KLD_RMS.history['val_acc']
      loss = history_KLD_RMS.history['loss']
      val_loss = history_KLD_RMS.history['val_loss']
      epochs = range(1, len(acc) + 1)
      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.legend()
      plt.figure()
      plt.plot(epochs, loss, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and validation loss')
      plt.legend()
      plt.show()
```



Avaliação da performance do modelo no conjunto de teste, utilizando o relatório de classificação. O relatório apresenta, para cada classe, as métricas precision, recall e F1-score, permitindo analisar detalhadamente os acertos e erros por classe.

```
[48]: y_true, y_pred = get_true_pred(modelS_KLD_RMS, test_dataset)
report = classification_report(y_true, y_pred, target_names=class_names,
                               output_dict=True)
class_only_report = {k: v for k, v in report.items() if k in class_names}
df = pd.DataFrame(class_only_report).T
print(df[['precision', 'recall', 'f1-score']].round(3))
```

	precision	recall	f1-score
buildings	0.856	0.762	0.806
forest	0.945	0.941	0.943
glacier	0.819	0.759	0.788
mountain	0.820	0.773	0.796
sea	0.757	0.924	0.832
street	0.862	0.876	0.869

```
2025-06-12 21:22:56.730716: I tensorflow/core/framework/local_rendevvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
```

6 Avaliação do melhor modelo

6.1 Comparação dos modelos utilizando a accuracy

```
[60]: val_loss_CatCross_RMS, val_acc_CatCross_RMS = modelS_CatCross_RMS.  
      ↪evaluate(validation_dataset)  
val_loss_CatCross_SGD, val_acc_CatCross_SGD = modelS_CatCross_SGD.  
      ↪evaluate(validation_dataset)  
val_loss_KLD_RMS, val_acc_KLD_RMS = modelS_KLD_RMS.evaluate(validation_dataset)  
val_loss_KLD_SGD, val_acc_KLD_SGD = modelS_KLD_SGD.evaluate(validation_dataset)  
  
print("Validation Accuracy dos modelos:")  
print(f"CatCross + RMSprop: {val_acc_CatCross_RMS:.4f}")  
print(f"CatCross + SGD      : {val_acc_CatCross_SGD:.4f}")  
print(f"KLD      + RMSprop: {val_acc_KLD_RMS:.4f}")  
print(f"KLD      + SGD      : {val_acc_KLD_SGD:.4f}")  
  
results = {  
    'CatCross_RMS': val_acc_CatCross_RMS,  
    'CatCross_SGD': val_acc_CatCross_SGD,  
    'KLD_RMS': val_acc_KLD_RMS,  
    'KLD_SGD': val_acc_KLD_SGD  
}  
  
# Identificação do melhor modelo com base na maior val_accuracy  
best_model = max(results, key=results.get)  
best_accuracy = results[best_model]  
  
print(f"\nMelhor modelo: {best_model} com val_accuracy = {best_accuracy:.4f}")
```

94/94 2s 16ms/step - acc:

0.8302 - loss: 0.4853

94/94 2s 19ms/step - acc:

0.7917 - loss: 0.5884

94/94 2s 18ms/step - acc:

0.8339 - loss: 0.4576

94/94 5s 52ms/step - acc:

0.7844 - loss: 0.5841

Validation Accuracy dos modelos:

CatCross + RMSprop: 0.8250

CatCross + SGD : 0.7920

KLD + RMSprop: 0.8190


KLD + SGD : 0.7810

Melhor modelo: CatCross_RMS com val_accuracy = 0.8250

6.2 Matriz de confusão do melhor modelo

```
[52]: y_true, y_pred = get_true_pred(modelS_CatCross_RMS, test_dataset)
      cm = confusion_matrix(y_true, y_pred)

      plt.figure(figsize=(6, 6))
      sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                  xticklabels=class_names,
                  yticklabels=class_names)
      plt.xlabel('Predicted')
      plt.ylabel('True')
      plt.title('Confusion Matrix - Melhor Modelo')
      plt.tight_layout()
      plt.show()
```



6.3 Calcular saída do modelo para uma imagem

```
[61]: img_path = 'Dataset/archive/seg_test/sea/20072.jpg'

img = tf.keras.preprocessing.image.load_img(
    img_path,
    target_size=(150, 150),
    interpolation='bilinear'
)

plt.imshow(img)
plt.axis('off')
plt.title("Imagem de Teste")
plt.show()

img_array = tf.keras.preprocessing.image.img_to_array(img)
img_array = tf.expand_dims(img_array, 0)

result = modelS_CatCross_RMS.predict(img_array)

class_names = ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
print("Probabilidades por classe:")
for i, prob in enumerate(result[0]):
    print(f"{class_names[i]:>10s}: {prob:.4f}")

predicted_class = np.argmax(result)
print(f"\nClasse prevista: {class_names[predicted_class]}")
↪({result[0][predicted_class]:.4f})"
```

Imagem de Teste



1/1 0s 120ms/step

Probabilidades por classe:

 buildings: 0.1919

 forest: 0.0034

 glacier: 0.0611

 mountain: 0.0504

 sea: 0.6622

 street: 0.0311

Classe prevista: sea (0.6622)