

modelS_2regularization

June 12, 2025

```
[9]: from tensorflow import keras
from keras import layers
from keras.preprocessing import image_dataset_from_directory
import matplotlib.pyplot as plt
from keras.utils import to_categorical
import tensorflow as tf
import numpy as np
from keras.preprocessing import image
from sklearn.metrics import classification_report
import seaborn as sns
import pandas as pd
from sklearn.metrics import confusion_matrix
import os, shutil
```

1 Funções

```
[10]: def get_true_pred(model, dataset):
    y_true = []
    y_pred = []
    for images, labels in dataset.unbatch().batch(1):
        y_true.append(np.argmax(labels.numpy()))
        pred = model.predict(images, verbose=0)
        y_pred.append(np.argmax(pred))
    return np.array(y_true), np.array(y_pred)
```

1.1 Carregamento do dataset

Carrega o dataset distribuido pelos diferentes conjuntos de dados.

```
[11]: train_dir = 'Dataset/archive/seg_train'
validation_dir = 'Dataset/archive/seg_val'
test_dir = 'Dataset/archive/seg_test'

train_buildings_dir = 'Dataset/archive/seg_train/buildings/'
train_forest_dir = 'Dataset/archive/seg_train/forest'
train_glacier_dir = 'Dataset/archive/seg_train/glacier'
train_mountain_dir = 'Dataset/archive/seg_train/mountain'
```

```

train_sea_dir = 'Dataset/archive/seg_train/sea'
train_street_dir = 'Dataset/archive/seg_train/street'

val_buildings_dir = 'Dataset/archive/seg_val/buildings'
val_forest_dir = 'Dataset/archive/seg_val/forest'
val_glacier_dir = 'Dataset/archive/seg_val/glacier'
val_mountain_dir = 'Dataset/archive/seg_val/mountain'
val_sea_dir = 'Dataset/archive/seg_val/sea'
val_street_dir = 'Dataset/archive/seg_val/street'

test_buildings_dir = 'Dataset/archive/seg_test/buildings'
test_forest_dir = 'Dataset/archive/seg_test/forest'
test_glacier_dir = 'Dataset/archive/seg_test/glacier'
test_mountain_dir = 'Dataset/archive/seg_test/mountain'
test_sea_dir = 'Dataset/archive/seg_test/sea'
test_street_dir = 'Dataset/archive/seg_test/street'

print('total training buildings images:', len(os.listdir(train_buildings_dir)))
print('total training forest images:', len(os.listdir(train_forest_dir)))
print('total training glacier images:', len(os.listdir(train_glacier_dir)))
print('total training mountain images:', len(os.listdir(train_mountain_dir)))
print('total training sea images:', len(os.listdir(train_sea_dir)))
print('total training street images:', len(os.listdir(train_street_dir)))

print('total validation buildings images:', len(os.listdir(val_buildings_dir)))
print('total validation forest images:', len(os.listdir(val_forest_dir)))
print('total validation glacier images:', len(os.listdir(val_glacier_dir)))
print('total validation mountain images:', len(os.listdir(val_mountain_dir)))
print('total validation sea images:', len(os.listdir(val_sea_dir)))
print('total validation street images:', len(os.listdir(val_street_dir)))

print('total test buildings images:', len(os.listdir(test_buildings_dir)))
print('total test forest images:', len(os.listdir(test_forest_dir)))
print('total test glacier images:', len(os.listdir(test_glacier_dir)))
print('total test mountain images:', len(os.listdir(test_mountain_dir)))
print('total test sea images:', len(os.listdir(test_sea_dir)))
print('total test street images:', len(os.listdir(test_street_dir)))

```

```

total training buildings images: 1691
total training forest images: 1771
total training glacier images: 1904
total training mountain images: 2012
total training sea images: 1774
total training street images: 1882
total validation buildings images: 500
total validation forest images: 500
total validation glacier images: 500
total validation mountain images: 500

```

```
total validation sea images: 500
total validation street images: 500
total test buildings images: 437
total test forest images: 474
total test glacier images: 553
total test mountain images: 525
total test sea images: 510
total test street images: 501
```

1.2 Distribuição de imagens por classe e por conjunto de dados

As imagens estão distribuídas por 3 conjuntos de dados: train, validation e test. Cada um desses conjuntos está distribuído por 6 classes: buildings, forest, glacier, mountain, sea e street.

1.2.1 Número total de imagens por classe:

Classe	Treino	Validação	Teste	Total
Buildings	1691	500	437	2628
Forest	1771	500	474	2745
Glacier	1904	500	553	2957
Mountain	2012	500	525	3037
Sea	1774	500	510	2784
Street	1882	500	501	2883
Total	11034	3000	3000	17034

1.2.2 Número total de imagens por conjunto de dados:

Conjunto de dados	Total
Treino	11034
Validação	3000
Teste	3000
Total geral	17034

2 Processamento dos dados

Carrega, redimensiona e organiza imagens em batches com rótulos one-hot, preparando os dados de treino, validação e teste.

```
[12]: IMG_SIZE = 150
      BATCH_SIZE = 32

      train_dataset = image_dataset_from_directory(
          train_dir,
          label_mode='categorical',
          image_size=(IMG_SIZE, IMG_SIZE),
          batch_size=BATCH_SIZE)
```

```

validation_dataset = image_dataset_from_directory(
    validation_dir,
    label_mode='categorical',
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE)

test_dataset = image_dataset_from_directory(
    test_dir,
    label_mode='categorical',
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE)

print(test_dataset)
class_names = train_dataset.class_names
print("Classes:", class_names)

```

```

Found 11034 files belonging to 6 classes.
Found 3000 files belonging to 6 classes.
Found 3000 files belonging to 6 classes.
<_PrefetchDataset element_spec=(TensorSpec(shape=(None, 150, 150, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None, 6), dtype=tf.float32,
name=None))>
Classes: ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']

```

3 Modelo (regularization: Dropout)

3.1 Criação da CNN

Criação da CNN que irá receber imagens de 150x150 píxeis, aplica normalização e passa por quatro camadas convolucionais com max pooling para extrair características, integrando camadas Dropout para reduzir overfitting (desligando aleatoriamente 50% dos neurónios durante o treino, como forma de regularização). A rede termina com uma camada densa com 512 unidades e uma camada de saída softmax para classificação multiclasse.

```

[13]: inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(32, 3, activation="relu")(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(64, 3, activation="relu")(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(128, 3, activation="relu")(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(128, 3, activation="relu")(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
x = layers.Dense(512, activation="relu")(x)

```

```

outputs = layers.Dense(len(class_names), activation="softmax")(x)
model_dropout = keras.Model(inputs, outputs)

print(model_dropout.summary())

```

Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 150, 150, 3)	0
rescaling_1 (Rescaling)	(None, 150, 150, 3)	0
conv2d_4 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_5 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_6 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_6 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_7 (Conv2D)	(None, 15, 15, 128)	147,584
max_pooling2d_7 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dropout_1 (Dropout)	(None, 6272)	0
dense_2 (Dense)	(None, 512)	3,211,776
dense_3 (Dense)	(None, 6)	3,078

Total params: 3,455,686 (13.18 MB)

Trainable params: 3,455,686 (13.18 MB)

Non-trainable params: 0 (0.00 B)

None

3.2 Compilação da CNN

Compilação da CNN utilizando a loss **categorical_crossentropy** e o optimizer **RMSprop**.

```
[14]: model_dropout.compile(optimizer=tf.keras.optimizers.  
    ↪RMSprop(learning_rate=1e-4), loss='categorical_crossentropy',  
    ↪metrics=['accuracy'])
```

3.3 Definição do callback

Definição de um callback que guarda automaticamente o modelo com a menor perda (loss) de validação durante o treino.

```
[15]: checkpoint_filepath = 'modelS_CatCross_RMS_dropout.keras'  
model_checkpoint_callback = keras.callbacks.ModelCheckpoint(  
    filepath=checkpoint_filepath,  
    monitor='val_loss',  
    save_best_only=True)
```

3.4 Treino da CNN

Treino da CNN durante 50 épocas utilizando o dataset de validação e o callback para guardar o melhor modelo.

```
[16]: history_dropout = model_dropout.fit(  
    train_dataset,  
    epochs=50,  
    validation_data=validation_dataset,  
    callbacks=[model_checkpoint_callback])
```

Epoch 1/50

345/345 21s 45ms/step -

accuracy: 0.4426 - loss: 1.3753 - val_accuracy: 0.5703 - val_loss: 1.1022

Epoch 2/50

345/345 12s 36ms/step -

accuracy: 0.6253 - loss: 0.9863 - val_accuracy: 0.6890 - val_loss: 0.8312

Epoch 3/50

345/345 15s 44ms/step -

accuracy: 0.6762 - loss: 0.8629 - val_accuracy: 0.7187 - val_loss: 0.7614

Epoch 4/50

345/345 12s 36ms/step -

accuracy: 0.7157 - loss: 0.7780 - val_accuracy: 0.7403 - val_loss: 0.7058

Epoch 5/50

345/345 15s 44ms/step -

accuracy: 0.7359 - loss: 0.7181 - val_accuracy: 0.7410 - val_loss: 0.7101

Epoch 6/50

345/345 12s 35ms/step -

accuracy: 0.7618 - loss: 0.6633 - val_accuracy: 0.7117 - val_loss: 0.7699
 Epoch 7/50
 345/345 13s 37ms/step -
 accuracy: 0.7794 - loss: 0.6241 - val_accuracy: 0.7193 - val_loss: 0.7779
 Epoch 8/50
 345/345 20s 58ms/step -
 accuracy: 0.7910 - loss: 0.5869 - val_accuracy: 0.7310 - val_loss: 0.7266
 Epoch 9/50
 345/345 13s 37ms/step -
 accuracy: 0.8069 - loss: 0.5492 - val_accuracy: 0.7847 - val_loss: 0.5986
 Epoch 10/50
 345/345 18s 52ms/step -
 accuracy: 0.8184 - loss: 0.5187 - val_accuracy: 0.7503 - val_loss: 0.6917
 Epoch 11/50
 345/345 14s 40ms/step -
 accuracy: 0.8293 - loss: 0.4868 - val_accuracy: 0.8077 - val_loss: 0.5454
 Epoch 12/50
 345/345 17s 50ms/step -
 accuracy: 0.8367 - loss: 0.4730 - val_accuracy: 0.7993 - val_loss: 0.5538
 Epoch 13/50
 345/345 15s 43ms/step -
 accuracy: 0.8411 - loss: 0.4459 - val_accuracy: 0.8173 - val_loss: 0.5046
 Epoch 14/50
 345/345 13s 38ms/step -
 accuracy: 0.8471 - loss: 0.4289 - val_accuracy: 0.8203 - val_loss: 0.5058
 Epoch 15/50
 345/345 17s 49ms/step -
 accuracy: 0.8541 - loss: 0.4198 - val_accuracy: 0.8033 - val_loss: 0.5352
 Epoch 16/50
 345/345 12s 34ms/step -
 accuracy: 0.8612 - loss: 0.3930 - val_accuracy: 0.8200 - val_loss: 0.5121
 Epoch 17/50
 345/345 15s 43ms/step -
 accuracy: 0.8732 - loss: 0.3742 - val_accuracy: 0.8250 - val_loss: 0.5069
 Epoch 18/50
 345/345 12s 35ms/step -
 accuracy: 0.8760 - loss: 0.3588 - val_accuracy: 0.8357 - val_loss: 0.4656
 Epoch 19/50
 345/345 14s 41ms/step -
 accuracy: 0.8774 - loss: 0.3497 - val_accuracy: 0.8357 - val_loss: 0.4636
 Epoch 20/50
 345/345 15s 44ms/step -
 accuracy: 0.8821 - loss: 0.3356 - val_accuracy: 0.7933 - val_loss: 0.5976
 Epoch 21/50
 345/345 12s 35ms/step -
 accuracy: 0.8876 - loss: 0.3286 - val_accuracy: 0.8353 - val_loss: 0.4769
 Epoch 22/50
 345/345 15s 44ms/step -

accuracy: 0.8889 - loss: 0.3143 - val_accuracy: 0.8260 - val_loss: 0.5241
 Epoch 23/50
 345/345 15s 43ms/step -
 accuracy: 0.8968 - loss: 0.2975 - val_accuracy: 0.8267 - val_loss: 0.5161
 Epoch 24/50
 345/345 14s 39ms/step -
 accuracy: 0.8995 - loss: 0.2850 - val_accuracy: 0.8427 - val_loss: 0.4812
 Epoch 25/50
 345/345 18s 51ms/step -
 accuracy: 0.9032 - loss: 0.2735 - val_accuracy: 0.8503 - val_loss: 0.4568
 Epoch 26/50
 345/345 13s 39ms/step -
 accuracy: 0.9129 - loss: 0.2665 - val_accuracy: 0.8410 - val_loss: 0.4671
 Epoch 27/50
 345/345 15s 44ms/step -
 accuracy: 0.9153 - loss: 0.2445 - val_accuracy: 0.8380 - val_loss: 0.4799
 Epoch 28/50
 345/345 13s 38ms/step -
 accuracy: 0.9145 - loss: 0.2428 - val_accuracy: 0.8267 - val_loss: 0.5130
 Epoch 29/50
 345/345 16s 46ms/step -
 accuracy: 0.9143 - loss: 0.2387 - val_accuracy: 0.8253 - val_loss: 0.5518
 Epoch 30/50
 345/345 13s 38ms/step -
 accuracy: 0.9203 - loss: 0.2263 - val_accuracy: 0.8347 - val_loss: 0.5003
 Epoch 31/50
 345/345 12s 35ms/step -
 accuracy: 0.9274 - loss: 0.2115 - val_accuracy: 0.8477 - val_loss: 0.4689
 Epoch 32/50
 345/345 16s 48ms/step -
 accuracy: 0.9297 - loss: 0.2033 - val_accuracy: 0.8477 - val_loss: 0.4951
 Epoch 33/50
 345/345 12s 35ms/step -
 accuracy: 0.9365 - loss: 0.1894 - val_accuracy: 0.8413 - val_loss: 0.5072
 Epoch 34/50
 345/345 16s 47ms/step -
 accuracy: 0.9357 - loss: 0.1891 - val_accuracy: 0.8330 - val_loss: 0.5603
 Epoch 35/50
 345/345 12s 35ms/step -
 accuracy: 0.9397 - loss: 0.1789 - val_accuracy: 0.8460 - val_loss: 0.4890
 Epoch 36/50
 345/345 13s 39ms/step -
 accuracy: 0.9372 - loss: 0.1730 - val_accuracy: 0.8323 - val_loss: 0.5264
 Epoch 37/50
 345/345 18s 51ms/step -
 accuracy: 0.9423 - loss: 0.1659 - val_accuracy: 0.8413 - val_loss: 0.4951
 Epoch 38/50
 345/345 13s 39ms/step -


```

accuracy: 0.9448 - loss: 0.1592 - val_accuracy: 0.8497 - val_loss: 0.5041
Epoch 39/50
345/345          15s 44ms/step -
accuracy: 0.9468 - loss: 0.1484 - val_accuracy: 0.8517 - val_loss: 0.4858
Epoch 40/50
345/345          13s 39ms/step -
accuracy: 0.9467 - loss: 0.1506 - val_accuracy: 0.8487 - val_loss: 0.4911
Epoch 41/50
345/345          13s 39ms/step -
accuracy: 0.9526 - loss: 0.1405 - val_accuracy: 0.8467 - val_loss: 0.5276
Epoch 42/50
345/345          16s 45ms/step -
accuracy: 0.9533 - loss: 0.1319 - val_accuracy: 0.8447 - val_loss: 0.5212
Epoch 43/50
345/345          12s 35ms/step -
accuracy: 0.9539 - loss: 0.1280 - val_accuracy: 0.8210 - val_loss: 0.7326
Epoch 44/50
345/345          15s 44ms/step -
accuracy: 0.9512 - loss: 0.1310 - val_accuracy: 0.8533 - val_loss: 0.4999
Epoch 45/50
345/345          12s 35ms/step -
accuracy: 0.9628 - loss: 0.1122 - val_accuracy: 0.8567 - val_loss: 0.5042
Epoch 46/50
345/345          12s 35ms/step -
accuracy: 0.9634 - loss: 0.1089 - val_accuracy: 0.8497 - val_loss: 0.5295
Epoch 47/50
345/345          15s 44ms/step -
accuracy: 0.9617 - loss: 0.1052 - val_accuracy: 0.8390 - val_loss: 0.6200
Epoch 48/50
345/345          13s 36ms/step -
accuracy: 0.9632 - loss: 0.1049 - val_accuracy: 0.8503 - val_loss: 0.5236
Epoch 49/50
345/345          17s 49ms/step -
accuracy: 0.9652 - loss: 0.1029 - val_accuracy: 0.8433 - val_loss: 0.5867
Epoch 50/50
345/345          12s 36ms/step -
accuracy: 0.9676 - loss: 0.0972 - val_accuracy: 0.8447 - val_loss: 0.5763

```

```

[17]: best_epoch = np.argmin(history_dropout.history['val_loss']) + 1
      print(f"Melhor época (menor val_loss): {best_epoch}")

```

Melhor época (menor val_loss): 25

3.5 Carregamento do modelo e validação


Carregamento e avaliação do modelo através do valor da accuracy.


```
[18]: model_dropout = keras.models.load_model('modelS_CatCross_RMS_dropout.keras')
      val_loss, val_acc = model_dropout.evaluate(validation_dataset)
      print('val_acc:', val_acc)
```

```
94/94          3s 21ms/step -
accuracy: 0.8406 - loss: 0.4978
val_acc: 0.85033333330154419
```

Representação gráfica dos valores da accuracy e da loss ao longo das épocas.

```
[19]: acc = history_dropout.history['accuracy']
      val_acc = history_dropout.history['val_accuracy']
      loss = history_dropout.history['loss']
      val_loss = history_dropout.history['val_loss']
      epochs = range(1, len(acc) + 1)
      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.legend()
      plt.figure()
      plt.plot(epochs, loss, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and validation loss')
      plt.legend()
      plt.show()
```





Avaliação da performance do modelo no conjunto de teste, utilizando o relatório de classificação. O relatório apresenta, para cada classe, as métricas precision, recall e F1-score, permitindo analisar detalhadamente os acertos e erros por classe.

```
[20]: y_true, y_pred = get_true_pred(model_dropout, test_dataset)
report = classification_report(y_true, y_pred, target_names=class_names,
                                output_dict=True)
class_only_report = {k: v for k, v in report.items() if k in class_names}
df = pd.DataFrame(class_only_report).T
print(df[['precision', 'recall', 'f1-score']].round(3))
```

```
2025-06-12 22:02:18.308742: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.12 = (f32[1,32,148,148]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,3,150,150]{3,2,1,0} %bitcast.262, f32[32,3,3,3]{3,2,1,0}
%bitcast.269, f32[32]{0} %bitcast.271), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1_1/conv2d_4_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
```

```

"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 22:02:18.371909: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.13 = (f32[1,64,72,72]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,32,74,74]{3,2,1,0} %bitcast.278, f32[64,32,3,3]{3,2,1,0}
%bitcast.285, f32[64]{0} %bitcast.287), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1_1/conv2d_5_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 22:02:18.425038: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.14 = (f32[1,128,34,34]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,64,36,36]{3,2,1,0} %bitcast.293, f32[128,64,3,3]{3,2,1,0}
%bitcast.300, f32[128]{0} %bitcast.302), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1_1/conv2d_6_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}
2025-06-12 22:02:18.482526: I
external/local_xla/xla/service/gpu/autotuning/conv_algorithm_picker.cc:549]
Omitted potentially buggy algorithm eng14{} for conv %cudnn-conv-bias-
activation.15 = (f32[1,128,15,15]{3,2,1,0}, u8[0]{0}) custom-
call(f32[1,128,17,17]{3,2,1,0} %bitcast.308, f32[128,128,3,3]{3,2,1,0}
%bitcast.315, f32[128]{0} %bitcast.317), window={size=3x3},
dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBiasActivationForward",
metadata={op_type="Conv2D" op_name="functional_1_1/conv2d_7_1/convolution"
source_file="/home/diogo/.pyenv/versions/3.10.18/lib/python3.10/site-
packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={
"operation_queue_id": "0", "wait_on_operation_queues": [], "cudnn_conv_backend_config": {"conv_result_scale": 1, "activation_mode": "kRelu", "side_input_scale": 0, "leakyrelu_alpha": 0}, "force_earliest_schedule": false}

```

	precision	recall	f1-score
buildings	0.864	0.796	0.829

forest	0.942	0.951	0.946
glacier	0.832	0.821	0.826
mountain	0.886	0.758	0.817
sea	0.776	0.931	0.847
street	0.863	0.880	0.872

2025-06-12 22:05:59.173953: I tensorflow/core/framework/local_rendezvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

4 Modelo (regularization: L2)

4.1 Criação da CNN

Criação da CNN que irá receber imagens de 150x150 píxeis, aplica normalização e passa por quatro camadas convolucionais com max pooling para extrair características, integrando regularização L2 nas camadas convolucionais para penalizar pesos excessivamente elevados e reduzir o overfitting. A rede termina com uma camada densa com 512 unidades e uma camada de saída softmax para classificação multiclasse.

```
[21]: inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(32, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(64, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(128, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(128, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Flatten()(x)
x = layers.Dense(512, activation="relu", kernel_regularizer='l2')(x)
outputs = layers.Dense(len(class_names), activation="softmax")(x)
model_l2 = keras.Model(inputs, outputs)

print(model_l2.summary())
```

Model: "functional_2"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 150, 150, 3)	0
rescaling_2 (Rescaling)	(None, 150, 150, 3)	0
conv2d_8 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_8 (MaxPooling2D)	(None, 74, 74, 32)	0

conv2d_9 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_9 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_10 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_10 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_11 (Conv2D)	(None, 15, 15, 128)	147,584
max_pooling2d_11 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_2 (Flatten)	(None, 6272)	0
dense_4 (Dense)	(None, 512)	3,211,776
dense_5 (Dense)	(None, 6)	3,078

Total params: 3,455,686 (13.18 MB)

Trainable params: 3,455,686 (13.18 MB)

Non-trainable params: 0 (0.00 B)

None

4.2 Compilação da CNN

Compilação da CNN utilizando a loss **categorical_crossentropy** e o optimizer **RMSprop**.

```
[22]: model_12.compile(
      loss='categorical_crossentropy',
      optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4),
      metrics=['acc'])
```

4.3 Definição do callback

Definição de um callback que guarda automaticamente o modelo com a menor perda (loss) de validação durante o treino.

```
[23]: checkpoint_filepath = 'modelS_CatCross_RMS_L2.keras'
      model_checkpoint_callback = keras.callbacks.ModelCheckpoint(
          filepath=checkpoint_filepath,
          monitor='val_loss',
```

```
save_best_only=True)
```

4.4 Treino da CNN

Treino da CNN durante 50 épocas utilizando o dataset de validação e o callback para guardar o melhor modelo.

```
[24]: history_12 = model_12.fit(  
    train_dataset,  
    epochs=50,  
    validation_data=validation_dataset,  
    callbacks=[model_checkpoint_callback])
```

Epoch 1/50

345/345 24s 57ms/step -

acc: 0.4528 - loss: 9.2746 - val_acc: 0.5573 - val_loss: 3.6021

Epoch 2/50

345/345 14s 40ms/step -

acc: 0.6091 - loss: 3.0461 - val_acc: 0.6103 - val_loss: 2.2840

Epoch 3/50

345/345 15s 44ms/step -

acc: 0.6483 - loss: 2.1270 - val_acc: 0.6237 - val_loss: 1.8955

Epoch 4/50

345/345 13s 39ms/step -

acc: 0.6690 - loss: 1.7784 - val_acc: 0.6527 - val_loss: 1.6544

Epoch 5/50

345/345 15s 45ms/step -

acc: 0.6777 - loss: 1.5969 - val_acc: 0.6843 - val_loss: 1.5193

Epoch 6/50

345/345 13s 38ms/step -

acc: 0.6881 - loss: 1.4691 - val_acc: 0.6867 - val_loss: 1.4160

Epoch 7/50

345/345 13s 36ms/step -

acc: 0.6998 - loss: 1.3742 - val_acc: 0.6973 - val_loss: 1.3263

Epoch 8/50

345/345 16s 46ms/step -

acc: 0.7055 - loss: 1.3055 - val_acc: 0.6417 - val_loss: 1.3917

Epoch 9/50

345/345 15s 43ms/step -

acc: 0.7095 - loss: 1.2612 - val_acc: 0.6767 - val_loss: 1.3005

Epoch 10/50

345/345 16s 48ms/step -

acc: 0.7205 - loss: 1.2088 - val_acc: 0.6947 - val_loss: 1.2515

Epoch 11/50

345/345 13s 38ms/step -

acc: 0.7268 - loss: 1.1717 - val_acc: 0.6557 - val_loss: 1.3328

Epoch 12/50

345/345 12s 36ms/step -

acc: 0.7234 - loss: 1.1442 - val_acc: 0.7117 - val_loss: 1.1288
 Epoch 13/50
 345/345 16s 45ms/step -
 acc: 0.7330 - loss: 1.1078 - val_acc: 0.7387 - val_loss: 1.1108
 Epoch 14/50
 345/345 12s 36ms/step -
 acc: 0.7431 - loss: 1.0844 - val_acc: 0.7290 - val_loss: 1.0947
 Epoch 15/50
 345/345 15s 45ms/step -
 acc: 0.7467 - loss: 1.0525 - val_acc: 0.7427 - val_loss: 1.0504
 Epoch 16/50
 345/345 12s 36ms/step -
 acc: 0.7504 - loss: 1.0380 - val_acc: 0.7607 - val_loss: 1.0250
 Epoch 17/50
 345/345 12s 35ms/step -
 acc: 0.7539 - loss: 1.0255 - val_acc: 0.7500 - val_loss: 1.0428
 Epoch 18/50
 345/345 15s 44ms/step -
 acc: 0.7574 - loss: 1.0046 - val_acc: 0.7377 - val_loss: 1.0705
 Epoch 19/50
 345/345 13s 37ms/step -
 acc: 0.7621 - loss: 0.9876 - val_acc: 0.7437 - val_loss: 1.0381
 Epoch 20/50
 345/345 17s 49ms/step -
 acc: 0.7615 - loss: 0.9803 - val_acc: 0.7690 - val_loss: 0.9599
 Epoch 21/50
 345/345 15s 44ms/step -
 acc: 0.7686 - loss: 0.9601 - val_acc: 0.6663 - val_loss: 1.1906
 Epoch 22/50
 345/345 13s 37ms/step -
 acc: 0.7692 - loss: 0.9513 - val_acc: 0.7447 - val_loss: 0.9894
 Epoch 23/50
 345/345 17s 48ms/step -
 acc: 0.7700 - loss: 0.9360 - val_acc: 0.7713 - val_loss: 0.9422
 Epoch 24/50
 345/345 14s 39ms/step -
 acc: 0.7770 - loss: 0.9198 - val_acc: 0.7597 - val_loss: 0.9706
 Epoch 25/50
 345/345 17s 48ms/step -
 acc: 0.7735 - loss: 0.9177 - val_acc: 0.7570 - val_loss: 0.9797
 Epoch 26/50
 345/345 13s 38ms/step -
 acc: 0.7800 - loss: 0.9060 - val_acc: 0.7603 - val_loss: 0.9483
 Epoch 27/50
 345/345 16s 46ms/step -
 acc: 0.7804 - loss: 0.8963 - val_acc: 0.7573 - val_loss: 0.9666
 Epoch 28/50
 345/345 13s 38ms/step -

acc: 0.7846 - loss: 0.8868 - val_acc: 0.7767 - val_loss: 0.9004
 Epoch 29/50
 345/345 13s 38ms/step -
 acc: 0.7788 - loss: 0.8784 - val_acc: 0.7787 - val_loss: 0.9022
 Epoch 30/50
 345/345 16s 47ms/step -
 acc: 0.7855 - loss: 0.8691 - val_acc: 0.7713 - val_loss: 0.9269
 Epoch 31/50
 345/345 13s 38ms/step -
 acc: 0.7904 - loss: 0.8646 - val_acc: 0.7317 - val_loss: 1.0009
 Epoch 32/50
 345/345 16s 47ms/step -
 acc: 0.7848 - loss: 0.8606 - val_acc: 0.7390 - val_loss: 0.9471
 Epoch 33/50
 345/345 13s 39ms/step -
 acc: 0.7880 - loss: 0.8504 - val_acc: 0.7880 - val_loss: 0.8730
 Epoch 34/50
 345/345 15s 43ms/step -
 acc: 0.7976 - loss: 0.8468 - val_acc: 0.7977 - val_loss: 0.8554
 Epoch 35/50
 345/345 16s 46ms/step -
 acc: 0.8033 - loss: 0.8367 - val_acc: 0.7833 - val_loss: 0.8614
 Epoch 36/50
 345/345 13s 38ms/step -
 acc: 0.7959 - loss: 0.8393 - val_acc: 0.7850 - val_loss: 0.8654
 Epoch 37/50
 345/345 15s 43ms/step -
 acc: 0.8020 - loss: 0.8177 - val_acc: 0.7797 - val_loss: 0.8835
 Epoch 38/50
 345/345 12s 36ms/step -
 acc: 0.8020 - loss: 0.8264 - val_acc: 0.7907 - val_loss: 0.8437
 Epoch 39/50
 345/345 13s 38ms/step -
 acc: 0.8036 - loss: 0.8110 - val_acc: 0.7697 - val_loss: 0.8976
 Epoch 40/50
 345/345 17s 49ms/step -
 acc: 0.8031 - loss: 0.8145 - val_acc: 0.7913 - val_loss: 0.8512
 Epoch 41/50
 345/345 13s 39ms/step -
 acc: 0.8071 - loss: 0.7967 - val_acc: 0.7913 - val_loss: 0.8296
 Epoch 42/50
 345/345 15s 44ms/step -
 acc: 0.8074 - loss: 0.8000 - val_acc: 0.7937 - val_loss: 0.8305
 Epoch 43/50
 345/345 12s 36ms/step -
 acc: 0.8044 - loss: 0.7901 - val_acc: 0.7920 - val_loss: 0.8220
 Epoch 44/50
 345/345 12s 35ms/step -

```

acc: 0.8097 - loss: 0.7814 - val_acc: 0.7517 - val_loss: 0.9344
Epoch 45/50
345/345          18s 51ms/step -
acc: 0.8110 - loss: 0.7858 - val_acc: 0.7743 - val_loss: 0.8539
Epoch 46/50
345/345          14s 40ms/step -
acc: 0.8159 - loss: 0.7818 - val_acc: 0.8013 - val_loss: 0.8134
Epoch 47/50
345/345          15s 44ms/step -
acc: 0.8148 - loss: 0.7744 - val_acc: 0.7523 - val_loss: 0.9336
Epoch 48/50
345/345          12s 36ms/step -
acc: 0.8151 - loss: 0.7692 - val_acc: 0.7953 - val_loss: 0.8094
Epoch 49/50
345/345          12s 36ms/step -
acc: 0.8191 - loss: 0.7626 - val_acc: 0.7987 - val_loss: 0.8033
Epoch 50/50
345/345          15s 44ms/step -
acc: 0.8134 - loss: 0.7701 - val_acc: 0.7700 - val_loss: 0.8718

```

```

[25]: best_epoch = np.argmin(history_l2.history['val_loss']) + 1
      print(f"Melhor época (menor val_loss): {best_epoch}")

```

Melhor época (menor val_loss): 49

4.5 Carregamento do modelo e validação

Carregamento e avaliação do modelo através do valor da accuracy.

```

[26]: model_l2 = keras.models.load_model('modelS_CatCross_RMS_L2.keras')
      val_loss, val_acc = model_l2.evaluate(validation_dataset)
      print('val_acc:', val_acc)

```

```

94/94          3s 21ms/step - acc:
0.7978 - loss: 0.8087
val_acc: 0.7986666560173035

```


Representação gráfica dos valores da accuracy e da loss ao longo das épocas.


```

[27]: acc = history_l2.history['acc']
      val_acc = history_l2.history['val_acc']
      loss = history_l2.history['loss']
      val_loss = history_l2.history['val_loss']
      epochs = range(1, len(acc) + 1)
      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.legend()
      plt.figure()

```

```
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```





Avaliação da performance do modelo no conjunto de teste, utilizando o relatório de classificação. O relatório apresenta, para cada classe, as métricas precision, recall e F1-score, permitindo analisar detalhadamente os acertos e erros por classe.

```
[28]: y_true, y_pred = get_true_pred(model_l2, test_dataset)
report = classification_report(y_true, y_pred, target_names=class_names,
                               output_dict=True)
class_only_report = {k: v for k, v in report.items() if k in class_names}
df = pd.DataFrame(class_only_report).T
print(df[['precision', 'recall', 'f1-score']].round(3))
```

	precision	recall	f1-score
buildings	0.789	0.787	0.788
forest	0.946	0.888	0.916
glacier	0.803	0.754	0.778
mountain	0.803	0.712	0.755
sea	0.748	0.843	0.793
street	0.778	0.868	0.821

```
2025-06-12 22:22:18.626637: I tensorflow/core/framework/local_rendezvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
```

5 Modelo (regularization: Dropout e L2)

5.1 Criação da CNN

Criação da CNN que irá receber imagens de 150x150 píxeis, aplica normalização e passa por quatro camadas convolucionais com max pooling para extrair características. Integra regularização L2 nas camadas convolucionais para penalizar pesos elevados e reduzir o overfitting, e camadas Dropout que desligam aleatoriamente 50% dos neurónios durante o treino. Termina com uma camada densa com 512 unidades e uma camada de saída softmax para classificação multiclasse.

```
[29]: inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(32, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(64, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(128, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Conv2D(128, 3, activation="relu", kernel_regularizer='l2')(x)
x = layers.MaxPooling2D(2)(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
x = layers.Dense(512, activation="relu", kernel_regularizer='l2')(x)
outputs = layers.Dense(len(class_names), activation="softmax")(x)
model_both = keras.Model(inputs, outputs)

print(model_both.summary())
```

Model: "functional_3"

Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 150, 150, 3)	0
rescaling_3 (Rescaling)	(None, 150, 150, 3)	0
conv2d_12 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_12 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_13 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_13 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_14 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_14 (MaxPooling2D)	(None, 17, 17, 128)	0

conv2d_15 (Conv2D)	(None, 15, 15, 128)	147,584
max_pooling2d_15 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_3 (Flatten)	(None, 6272)	0
dropout_2 (Dropout)	(None, 6272)	0
dense_6 (Dense)	(None, 512)	3,211,776
dense_7 (Dense)	(None, 6)	3,078

Total params: 3,455,686 (13.18 MB)

Trainable params: 3,455,686 (13.18 MB)

Non-trainable params: 0 (0.00 B)

None

5.2 Compilação da CNN

Compilação da CNN utilizando a loss **categorical_crossentropy** e o optimizer **RMSprop**.

```
[30]: model_both.compile(optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4),
↳ loss='categorical_crossentropy', metrics=['accuracy'])
```

5.3 Definição do callback

Definição de um callback que guarda automaticamente o modelo com a menor perda (loss) de validação durante o treino.

```
[31]: checkpoint_filepath = 'modelS_CatCross_RMS_dropout_L2.keras'
model_checkpoint_callback = keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_filepath,
    monitor='val_loss',
    save_best_only=True)
```

5.4 Treino da CNN

Treino da CNN durante 50 épocas utilizando o dataset de validação e o callback para guardar o melhor modelo.

```
[32]: history_both = model_both.fit(  
      train_dataset,  
      epochs=50,  
      validation_data=validation_dataset,  
      callbacks=[model_checkpoint_callback])
```

```
Epoch 1/50  
345/345          23s 55ms/step -  
accuracy: 0.4214 - loss: 9.5866 - val_accuracy: 0.5413 - val_loss: 3.9742  
Epoch 2/50  
345/345          12s 36ms/step -  
accuracy: 0.5933 - loss: 3.3751 - val_accuracy: 0.6030 - val_loss: 2.4343  
Epoch 3/50  
345/345          17s 50ms/step -  
accuracy: 0.6382 - loss: 2.2580 - val_accuracy: 0.6430 - val_loss: 1.9492  
Epoch 4/50  
345/345          14s 40ms/step -  
accuracy: 0.6604 - loss: 1.8401 - val_accuracy: 0.6463 - val_loss: 1.7115  
Epoch 5/50  
345/345          12s 36ms/step -  
accuracy: 0.6727 - loss: 1.6196 - val_accuracy: 0.6677 - val_loss: 1.5477  
Epoch 6/50  
345/345          15s 43ms/step -  
accuracy: 0.6845 - loss: 1.4849 - val_accuracy: 0.6790 - val_loss: 1.4433  
Epoch 7/50  
345/345          12s 34ms/step -  
accuracy: 0.6917 - loss: 1.4102 - val_accuracy: 0.6963 - val_loss: 1.3490  
Epoch 8/50  
345/345          15s 43ms/step -  
accuracy: 0.7043 - loss: 1.3289 - val_accuracy: 0.7130 - val_loss: 1.2866  
Epoch 9/50  
345/345          12s 34ms/step -  
accuracy: 0.7116 - loss: 1.2796 - val_accuracy: 0.7273 - val_loss: 1.2266  
Epoch 10/50  
345/345          12s 36ms/step -  
accuracy: 0.7154 - loss: 1.2340 - val_accuracy: 0.7120 - val_loss: 1.2114  
Epoch 11/50  
345/345          16s 47ms/step -  
accuracy: 0.7235 - loss: 1.2049 - val_accuracy: 0.7253 - val_loss: 1.1907  
Epoch 12/50  
345/345          13s 37ms/step -  
accuracy: 0.7254 - loss: 1.1677 - val_accuracy: 0.7343 - val_loss: 1.1421  
Epoch 13/50  
345/345          13s 37ms/step -  
accuracy: 0.7339 - loss: 1.1366 - val_accuracy: 0.7480 - val_loss: 1.1096  
Epoch 14/50  
345/345          15s 44ms/step -  
accuracy: 0.7413 - loss: 1.1195 - val_accuracy: 0.7410 - val_loss: 1.0936
```


Epoch 15/50
 345/345 12s 36ms/step -
 accuracy: 0.7451 - loss: 1.0907 - val_accuracy: 0.7560 - val_loss: 1.0573

Epoch 16/50
 345/345 15s 44ms/step -
 accuracy: 0.7547 - loss: 1.0684 - val_accuracy: 0.7287 - val_loss: 1.1081

Epoch 17/50
 345/345 13s 36ms/step -
 accuracy: 0.7589 - loss: 1.0490 - val_accuracy: 0.7550 - val_loss: 1.0485

Epoch 18/50
 345/345 12s 35ms/step -
 accuracy: 0.7567 - loss: 1.0374 - val_accuracy: 0.7027 - val_loss: 1.1904

Epoch 19/50
 345/345 15s 45ms/step -
 accuracy: 0.7523 - loss: 1.0343 - val_accuracy: 0.7653 - val_loss: 1.0116

Epoch 20/50
 345/345 12s 35ms/step -
 accuracy: 0.7671 - loss: 1.0063 - val_accuracy: 0.7590 - val_loss: 1.0176

Epoch 21/50
 345/345 15s 44ms/step -
 accuracy: 0.7688 - loss: 1.0036 - val_accuracy: 0.7453 - val_loss: 1.0283

Epoch 22/50
 345/345 12s 36ms/step -
 accuracy: 0.7721 - loss: 0.9823 - val_accuracy: 0.7770 - val_loss: 0.9633

Epoch 23/50
 345/345 12s 35ms/step -
 accuracy: 0.7773 - loss: 0.9660 - val_accuracy: 0.7680 - val_loss: 0.9873

Epoch 24/50
 345/345 15s 44ms/step -
 accuracy: 0.7855 - loss: 0.9553 - val_accuracy: 0.7550 - val_loss: 1.0057

Epoch 25/50
 345/345 14s 41ms/step -
 accuracy: 0.7747 - loss: 0.9548 - val_accuracy: 0.7657 - val_loss: 0.9789

Epoch 26/50
 345/345 16s 47ms/step -
 accuracy: 0.7827 - loss: 0.9414 - val_accuracy: 0.7717 - val_loss: 0.9654

Epoch 27/50
 345/345 15s 43ms/step -
 accuracy: 0.7881 - loss: 0.9302 - val_accuracy: 0.7860 - val_loss: 0.9220

Epoch 28/50
 345/345 15s 42ms/step -
 accuracy: 0.7917 - loss: 0.9171 - val_accuracy: 0.7773 - val_loss: 0.9257

Epoch 29/50
 345/345 18s 52ms/step -
 accuracy: 0.7876 - loss: 0.9184 - val_accuracy: 0.7900 - val_loss: 0.9193

Epoch 30/50
 345/345 16s 46ms/step -
 accuracy: 0.7851 - loss: 0.9170 - val_accuracy: 0.7693 - val_loss: 0.9448

Epoch 31/50
 345/345 19s 56ms/step -
 accuracy: 0.7961 - loss: 0.8937 - val_accuracy: 0.7847 - val_loss: 0.9126

Epoch 32/50
 345/345 16s 46ms/step -
 accuracy: 0.8031 - loss: 0.8880 - val_accuracy: 0.7923 - val_loss: 0.8883

Epoch 33/50
 345/345 19s 55ms/step -
 accuracy: 0.7948 - loss: 0.8906 - val_accuracy: 0.7843 - val_loss: 0.9122

Epoch 34/50
 345/345 16s 47ms/step -
 accuracy: 0.7957 - loss: 0.8809 - val_accuracy: 0.8060 - val_loss: 0.8578

Epoch 35/50
 345/345 19s 55ms/step -
 accuracy: 0.8061 - loss: 0.8708 - val_accuracy: 0.7693 - val_loss: 0.9256

Epoch 36/50
 345/345 16s 48ms/step -
 accuracy: 0.7998 - loss: 0.8646 - val_accuracy: 0.7983 - val_loss: 0.8543

Epoch 37/50
 345/345 19s 55ms/step -
 accuracy: 0.8014 - loss: 0.8579 - val_accuracy: 0.7773 - val_loss: 0.8929

Epoch 38/50
 345/345 17s 50ms/step -
 accuracy: 0.7971 - loss: 0.8637 - val_accuracy: 0.7817 - val_loss: 0.8797

Epoch 39/50
 345/345 19s 54ms/step -
 accuracy: 0.8071 - loss: 0.8506 - val_accuracy: 0.7843 - val_loss: 0.8952

Epoch 40/50
 345/345 15s 45ms/step -
 accuracy: 0.8056 - loss: 0.8435 - val_accuracy: 0.7850 - val_loss: 0.8862

Epoch 41/50
 345/345 19s 53ms/step -
 accuracy: 0.8092 - loss: 0.8330 - val_accuracy: 0.7770 - val_loss: 0.8985

Epoch 42/50
 345/345 17s 49ms/step -
 accuracy: 0.8050 - loss: 0.8409 - val_accuracy: 0.7893 - val_loss: 0.8632

Epoch 43/50
 345/345 20s 57ms/step -
 accuracy: 0.8116 - loss: 0.8290 - val_accuracy: 0.7970 - val_loss: 0.8492

Epoch 44/50
 345/345 17s 50ms/step -
 accuracy: 0.8162 - loss: 0.8223 - val_accuracy: 0.7787 - val_loss: 0.9015

Epoch 45/50
 345/345 19s 54ms/step -
 accuracy: 0.8147 - loss: 0.8233 - val_accuracy: 0.7230 - val_loss: 1.0055

Epoch 46/50
 345/345 21s 60ms/step -
 accuracy: 0.8109 - loss: 0.8214 - val_accuracy: 0.7993 - val_loss: 0.8313

```

Epoch 47/50
345/345          16s 45ms/step -
accuracy: 0.8149 - loss: 0.8134 - val_accuracy: 0.8003 - val_loss: 0.8389
Epoch 48/50
345/345          20s 57ms/step -
accuracy: 0.8117 - loss: 0.8128 - val_accuracy: 0.7687 - val_loss: 0.9084
Epoch 49/50
345/345          17s 48ms/step -
accuracy: 0.8175 - loss: 0.8061 - val_accuracy: 0.8057 - val_loss: 0.8178
Epoch 50/50
345/345          22s 63ms/step -
accuracy: 0.8126 - loss: 0.7987 - val_accuracy: 0.7853 - val_loss: 0.8561

```

5.5 Carregamento do modelo e validação

Carregamento e avaliação do modelo através do valor da accuracy.

```

[34]: model_both = keras.models.load_model('modelS_CatCross_RMS_dropout_L2.keras')
      val_loss, val_acc = model_both.evaluate(validation_dataset)
      print('val_acc:', val_acc)

```

```

94/94          4s 24ms/step -
accuracy: 0.8075 - loss: 0.8293
val_acc: 0.8056666851043701


```


Representação gráfica dos valores da accuracy e da loss ao longo das épocas.

```

[36]: acc = history_both.history['accuracy']
      val_acc = history_both.history['val_accuracy']
      loss_CatCros = history_both.history['loss']
      val_loss = history_both.history['val_loss']
      epochs = range(1, len(acc) + 1)
      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.legend()
      plt.figure()
      plt.plot(epochs, loss_CatCros, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and validation loss')
      plt.legend()
      plt.show()

```





Avaliação da performance do modelo no conjunto de teste, utilizando o relatório de classificação. O relatório apresenta, para cada classe, as métricas precision, recall e F1-score, permitindo analisar detalhadamente os acertos e erros por classe.

```
[37]: y_true, y_pred = get_true_pred(model_both, test_dataset)
report = classification_report(y_true, y_pred, target_names=class_names,
                               output_dict=True)
class_only_report = {k: v for k, v in report.items() if k in class_names}
df = pd.DataFrame(class_only_report).T
print(df[['precision', 'recall', 'f1-score']].round(3))
```

	precision	recall	f1-score
buildings	0.820	0.721	0.767
forest	0.853	0.958	0.903
glacier	0.766	0.817	0.791
mountain	0.815	0.737	0.774
sea	0.786	0.829	0.807
street	0.840	0.806	0.823

6 Avaliação do melhor modelo

6.1 Comparação dos modelos utilizando a accuracy

```
[38]: val_loss_CatCross_RMS_dropout, val_acc_CatCross_RMS_dropout = model_dropout.  
      ↪evaluate(validation_dataset)  
val_loss_CatCross_RMS_l2, val_acc_CatCross_RMS_l2 = model_l2.  
      ↪evaluate(validation_dataset)  
val_loss_CatCross_RMS_dropout_l2, val_acc_CatCross_RMS_dropout_l2 = model_both.  
      ↪evaluate(validation_dataset)  
  
print("Validation Accuracy dos modelos:")  
print(f"CatCross + RMSprop + Dropout: {val_loss_CatCross_RMS_dropout:.4f}")  
print(f"CatCross + RMSprop + L2: {val_loss_CatCross_RMS_l2:.4f}")  
print(f"CatCross + RMSprop + Dropout + L2: {val_loss_CatCross_RMS_dropout_l2:.  
      ↪4f}")  
  
results = {  
    'CatCross_RMS_Dropout': val_acc_CatCross_RMS_dropout,  
    'CatCross_RMS_L2': val_acc_CatCross_RMS_l2,  
    'CatCross_RMS_Dropout_L2': val_acc_CatCross_RMS_dropout_l2  
}  
  
# Identificar o melhor modelo com base na maior val_accuracy  
best_model = max(results, key=results.get)  
best_accuracy = results[best_model]  
  
print(f"\nMelhor modelo: {best_model} com val_accuracy = {best_accuracy:.4f}")
```

```
94/94          4s 35ms/step -  
accuracy: 0.8360 - loss: 0.4948  
94/94          3s 30ms/step - acc:  
0.7997 - loss: 0.8191  
94/94          3s 29ms/step -  
accuracy: 0.8049 - loss: 0.8338  
Validation Accuracy dos modelos:  
CatCross + RMSprop + Dropout: 0.4568  
CatCross + RMSprop + L2: 0.8033  
CatCross + RMSprop + Dropout + L2: 0.8178
```

Melhor modelo: CatCross_RMS_Dropout com val_accuracy = 0.8503

6.2 Matriz de confusão do melhor modelo


```
[39]: y_true, y_pred = get_true_pred(model_dropout, test_dataset)  
cm = confusion_matrix(y_true, y_pred)  
  
plt.figure(figsize=(6, 6))
```

```

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names,
            yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix - Melhor Modelo')
plt.tight_layout()
plt.show()

```

2025-06-12 22:52:36.305303: I tensorflow/core/framework/local_rendezvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence



6.3 Calcular saída do modelo para uma imagem

```
[40]: img_path = 'Dataset/archive/seg_test/sea/20072.jpg'

img = tf.keras.preprocessing.image.load_img(
    img_path,
    target_size=(150, 150),
    interpolation='bilinear'
)

plt.imshow(img)
plt.axis('off')
plt.title("Imagem de Teste")
plt.show()

img_array = tf.keras.preprocessing.image.img_to_array(img)
img_array = tf.expand_dims(img_array, 0)

# Previsão
result = model_dropout.predict(img_array)

class_names = ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
print("Probabilidades por classe:")
for i, prob in enumerate(result[0]):
    print(f"{class_names[i]:>10s}: {prob:.4f}")

# Classe prevista
predicted_class = np.argmax(result)
print(f"\nClasse prevista: {class_names[predicted_class]}")
↪({result[0][predicted_class]:.4f})"
```


Imagem de Teste



1/1 0s 41ms/step

Probabilidades por classe:

 buildings: 0.0069

 forest: 0.0004

 glacier: 0.0563

 mountain: 0.0155

 sea: 0.9200

 street: 0.0009

Classe prevista: sea (0.9200)